

IPv6 Introduction

IPv6 core protocol definition is in RFC 2460

RIPng, EIGRPv6, OSPFv3, ICMPv6, traceroute6, ping6, MP BGP-4 (multiprotocol BGP v4)

ARP is replaced with NDP

Truncating addresses

4 digits informally a "quartet", so $4 \times 4 = 16$

0000:0000:0000:0000:0000:0000:0000:0000 = 128 bits

The Rules:

Leading 0's can be removed so 003A is 3A

Sets of consecutive 0's can be truncated to one in an address - 0000:0000 to 0:0

Replacing multiple quartets of 0's (0000:0000 with ::) can only be done once.

Doing the 3 in one address should be ok, within those rules

Expanding is easy: inverse the rules and pad out so there are 8 quartets of 4 bits each.

2000:1234:5678:9ABC::/64 prefix length is 64 bits (4 octets, 16 each)

In this way, a "/" prefix notation is almost like a CIDR - it tells how many bits used.

Given Address	Expanded Form	Address's Prefix
34BA:B:B:0:5555:0:6060:707/80	34BA:000B:000B:0000:5555:0000:6060:0707	34BA:B:B:0:5555::/80
3124::DEAD:CAFE:FF:FE00:1/80	3124:0000:0000:DEAD:CAFE:00FF:FE00:0001	3124:0:0:DEAD:CAFE::/80
2BCD::FACE:BEFF:FEFE:CAFE/48	2BCD:0000:0000:0000:FACE:BEFF:FEFE:CAFE	2BCD::/48
3FED:F:E0:D00:FACE:BAFF:FE00:0/48	3FED:000F:00E0:0D00:FACE:BAFF:FE00:0000	3FED:F:E0::/48
210F:A:B:C:CCCC:B0B0:9999:9009/40	210F:000A:000B:000C:CCCC:B0B0:9999:9009	210F:A::/40
34BA:B:B:0:5555:0:6060:707/36	34BA:000B:000B:0000:5555:0000:6060:0707	34BA:B::/36
3124::DEAD:CAFE:FF:FE00:1/60	3124:0000:0000:DEAD:CAFE:00FF:FE00:0001	3124:0:0:DEA0::/60
2BCD::FACE:1:BEFF:FEFE:CAFE/56	2BCD:0000:0000:FACE:0001:BEFF:FEFE:CAFE	2BCD:0:0:FA00::/56
3FED:F:E0:D000:FACE:BAFF:FE00:0/52	3FED:000F:00E0:D000:FACE:BAFF:FE00:0000	3FED:F:E0:D000::/52
3BED:800:0:40:FACE:BAFF:FE00:0/44	3BED:0800:0000:0040:FACE:BAFF:FE00:0000	3BED:800::/44

Global unicasts are initially delegated in blocks by IANA, ARIN etc., and are leased for usage by various service providers. Since there are so many IPv6 addresses, they exist as just a utility for network operation provisioning- there is no recognized range of addresses that is "owned" or has brand recognition (like Google's name server address of 8.8.8.8), you just get any block the size you need from an ISP and move it to another as needed.

2001:db8:2222::/48 -Company A

2001:db8:3333::/48 -Company B

And here are some global unicast versions of subnet and hostnames:

2001:db8:3333:0001::/64 -Company B's subnet 1

2001:db8:3333:0002::/64 -Company B's subnet 2

2001:db8:3333:2::1 -A host in Company B's subnet 2

2001:db8:3333:0003::/64 -Company B's subnet 3

2001:db8:3333:3::10 -A host in Company B's subnet 3

Initially started with 2 or 3; now anything not reserved

Global addresses have the same overall structure:

Global routing prefix (IANA- 48 bits) + subnet bits (16) + interfaceID (typically 64 bits) = 128 bits total

2001:0db8:1111 (48 bits): 0001(16 bits) :0000:0000:0000:0001 (64 bits)

In this example, subnets would be 0001, 0002, 0003

Unique Local Unicast Addresses are like private IPv4 addresses

FD (8 bits) - GlobalID (pseudorandom 40bits) - subnet 16 bits - remaining 64 bits for interface

You prefix with /48 and then add the subnet (64 total) and interface (the last 64 bits)

- FD00:1:1:0001::/64 would be a sample unique local

- FD + 00:1234:5678: + 9ABC:DEF1:2345:6789:ABCD

Static Unicast Address Configuration

Generating a Unique Interface ID using EUI-64 is usually done for us just by assigning the address with:
ipv6 address 2001:db8:1111:1::/64 eui-64

EUI-64 creates the interface ID part as follows:

Split the 6-byte (12-hex) MAC address in two halves, and insert FFFE in between the two

Invert the seventh bit of the interface ID. Examples:

Interface MAC address is aa12:bcbc:1234

10101010 represents the first 8 bits of the MAC address (aa) - when inverting the 7th bit it becomes 10101000.

The answer is A8 and the EUI-64 address: 2001:0db8:0:1:a812:bcbf:febc:1234 EUI-64

MAC address 0c0c:dede:1234

0c is 00001100 in the first 8 bits of the MAC address, which then becomes 00001110 when flipping the 7th bit.

The answer is then 0e and the EUI-64 address: 2001:0db8:0:1:0e0c:deff:fede:1234 EUI-64

MAC address 0b34:ba12:1234

0b in binary is 00001011, the first 8 bits of the MAC address, which then becomes 00001001..”

The answer is 09, and the IPv6 EUI-64 address: 2001:0db8:0:1:0934:baff:fe12:1234 EUI-64

```
interface GigabitEthernet0/0
```

```
  ipv6 address 2001:db8:1111:1::/64 eui-64
```

```
interface serial0/0/0
```

```
  ipv6 address 2001:db8:1111:2::/64 eui-64
```

```
show ipv6 interface brief
```

```
  GigabitEthernet0/0 [up/up]
```

```
    FE80::1FF:FE01:101
```

```
    <---link local, employing EUI64
```

```
    2001:DB8:1111:1:0:1FF:FE01:101
```

```
    <---EUI64 global unicast address
```

```
  Serial0/0 [up/up]
```

```
    FE80::1FF:FE01:101
```

```
    2001:DB8:1111:2:0:1FF:FE01:101
```

Link-Local Addresses (LLA)

Auto-generated - interfaces can create their own LLA

Common uses - Overhead protocols (NDP) and next-hop address

Forwarding scope is the local link only - (for within a subnet)

Unicast, represent a single host. Always starts with the same prefix: the first 10 bits are FE80::/10

64 Bits - FE80:0000:0000:0000 + 64 Bits - Interface ID: EUI-64

Second half can be formed with different rules

EUI-64 (like above)

Operating Systems use random processes

Manually configured

Key IPv6 Local-Scope Multicast Addresses

Short Name	Multicast Address	Meaning	IPv4 Equivalent
All-nodes	FF02::1	All nodes (using IPv6)	Subnet Broadcast Address
All-routers	FF02::2	All routers (using IPv6)	None
All-OSPF	FF02::5	All OSPF routers	224.0.0.5
All-OSPF-DR	FF02::6	All OSPF designated routers	224.0.0.6
EIGRPv6 routers	FF02::A	All routers (using EIGRPv6)	224.0.0.10

```
R1# show ipv6 interface GigabitEthernet 0/0
```

```
GigabitEthernet0/0 is up, line protocol is up
```

```
  IPv6 is enabled, link-local address is FE80::1FF:FE01:101
```

```
  [extra info omitted]
```

```
  Joined group address(es):
```

```
    FF02::1
```

```
    FF02::2
```

FF02::A
FF02::1:FF01:101 <-----This host's solicited node multicast

Solicited-Node Multicast Addresses

A multicast address that is link-local in scope
Based on the unicast of the host, and only the last six hex digits
Each host listens for packets sent to **its** solicited-node address

Format: [Defined by RFC] + [Last 6 Hex Digits of Unicast]

Example: FF02:0000:0000:0000:0000:0001:FF__:_:_:_

By design, all hosts with the same last six hex digits use the same address
[FF02::1:FF/108 ?]

Special addresses used by routers:

:: (all 0's) signify unknown/ unspecified address Used when a host's own address is not yet known

::1 loopback - 127 binary 0s with a single 1

Neighbor Discovery Protocol (NDP):

Router discovery - available routers in the same subnet

- Router solicitation (RS) - Over FF02::1 (all routers MC), asks routers to "sound off"
- Router advertisements (RA) - Over FF02::1 (all nodes MC), router replies with link-local address, other info

Neighbor MAC discovery. NDP replaces ARP

- Neighbor Solicitation (NS, like an ARP request) and Neighbor Advertisement (NA, like an ARP reply)
- Duplicated Address Detection (DAD) uses NS requests and NA replies to keep track of unicasts

(SLAAC uses NS/NA's in DAD to do DHCP autoconfig)

Dynamic Unicast Address Configuration

Routers can be configured to use dynamic addresses

Two methods: interface FastEthernet0/0

ipv6 address dhcp <-----Using Stateful DHCP

ipv6 address autoconfig <----- Using Stateless Address Autoconfiguration (SLAAC)

Regular DHCPv6

DHCP client sends out a solicitation using it's own link-local as the source address, and the "All-DHCP-Agents" MC - ff02::1:2) as it's destination

Configuring R1 to be a DHCP relay:

interface GigabitEthernet0/0

ipv6 dhcp relay destination 2001:db8:1111:3::8

R1# show ipv6 interface g0/0

GigabitEthernet0/0 is up, line protocol is up

Joined group address(es):

FF02::1

FF02::2

FF02::A

FF02::1:2 <----- **this got added for the relay role**

FF02::1:FF00:1

SLAAC: Using Stateless Address Autoconfiguration

Building an IPv6 Address Using SLAAC is easy:

A NDP RS is sent to get an RA to provide the local link prefix

The host then slaps on a EUI-64 or random value for an interface ID

Send out Neighbor Solicitation for DAD process to verify it's unique.

DHCPv6 vs SLAAC?

With stateless SLAAC, you get your DNS servers from the local router, and you don't need to lease an address.

With a DHCP server being involved, you can have a lease and state information on a client. You might think DHCP would give more control, but it really doesn't since users can still use SLAAC instead.

CHECK THIS - question in practice set says SLAAC gets NDP to give it prefix length, default router addresses, but that Stateless DHCP provides the DNS servers.

Ping6 and Traceroute6, all the same
Extended is still just "ping" since it asks you the protocol

/128 is a "host route" for the router ip6 address

show ipv6 neighbors
show ipv6 routers -- to show connected routers

sh ipv6 interface GigabitEthernet 0/0

show ipv6 interface brief <--- no prefix length info
show ipv6 interface <--- Full interface details

ipv6 route ::/0
sh ipv6 neigh

Configuring static IP6 on two routers:

```
!R1
ipv6 unicast-routing      <---If not enabled, the router will act like a host and won't route
interface gigabitethernet0/0
  ipv6 address 2001:Ddb8:1111:1::1/64
interface serial0/0/0
  ipv6 address 2001:0db8:1111:0002:0000:0000:0000:0001/64
!R2
ipv6 unicast-routing
interface gigabitethernet0/0
  ipv6 address 2001:db8:1111:3::2/64
interface serial0/0/1
  ipv6 address 2001:db8:1111:2::2/64
```

#R1 show ipv6 route connected

```
IPv6 Routing Table - default 5 entries
!Omitted
C 2001:DB8:1111:1::/64 [0/0]
  via GigabitEthernet0/0, directly connected
C 2001:DB8:1111:2::/64 [0/0]
  via Serial0/0/0, directly connected
```

#R1 show ipv6 interface GigabitEthernet 0/0

```
GigabitEthernet0/0 is up, line protocol is up
IPv6 is enabled, link-local address is FE80:1FF:FE01:101
No virtual link-local address(es):
Description: LAN at Site 1
Global unicast address(es):
```

2001:DB8:1111:1::1, subnet is 2001:DB8:1111:1::/64
#R1 **show ipv6 interface s0/0/0**
Serial0/0/0 is up, line protocol is up
IPv6 is enabled, link-local address is FE80:1FF:FE01:101
No virtual link-local address(es):
Description: Link to R2
Global unicast address(es):
2001:DB8:1111:2::1, subnet is 2001:DB8:1111:2::/64

#R1 **show ipv6 interface brief**
GigabitEthernet0/0 [up/up]
FE80:1FF:FE01:101
2001:DB8:1111:1::1
GigabitEthernet0/1 [administratively down/down]
unassigned
Serial0/0 [up/up]
FE80:1FF:FE01:101
2001:DB8:1111:2::1

<http://freeccnalab.com/>
<http://www.packettracelab.com/>
<http://www.packettracernetwork.com/labs/packettracerlabs.html>
<https://boubakr92.wordpress.com/2013/09/16/ccna-cheat-sheet-part-1/>

ipv6 address <command>
ipv6 address
sh ipv6 route

- **Adding IPv6 Routes to Routing Tables**
 - Configuration of IPv6 addresses on working interfaces
 - Direct configuration of a static route
 - Configuration of a routing protocol on routers that share the same link
- **Rules for Connected and Local Routes**
 - Routers create routes based on each unicast address on an interface
 - Router creates a route for the subnet (a connected route)
 - Router creates a host route (/128) for the router IPv6 address
 - Routers do not create routes based on the LLA for the interface
 - Routers remove the connected and local routes if the interface fails
 - Re-adds when the interface is up/up

ipv6 unicast-routing

!

interface serial0/0/0

ipv6 address 2001:db8:1111:4::1/64

!

interface serial0/0/1

ipv6 address 2001:db8:1111:5::1/64

!

interface gigabitethernet0/0

ipv6 address 2001:db8:1111:1::1/64

Before adding static or OSPF routes:

show ipv6 route

!Omitted

Codes: C - connected, L - local....

!Omitted

C 2001:db8:1111:1::/64 [0/0]

via GigabitEthernet0/0, directly connected

!Omitted

C 2001:db8:1111:4::/64 [0/0]

via Serial0/0/0, directly connected

!Omitted

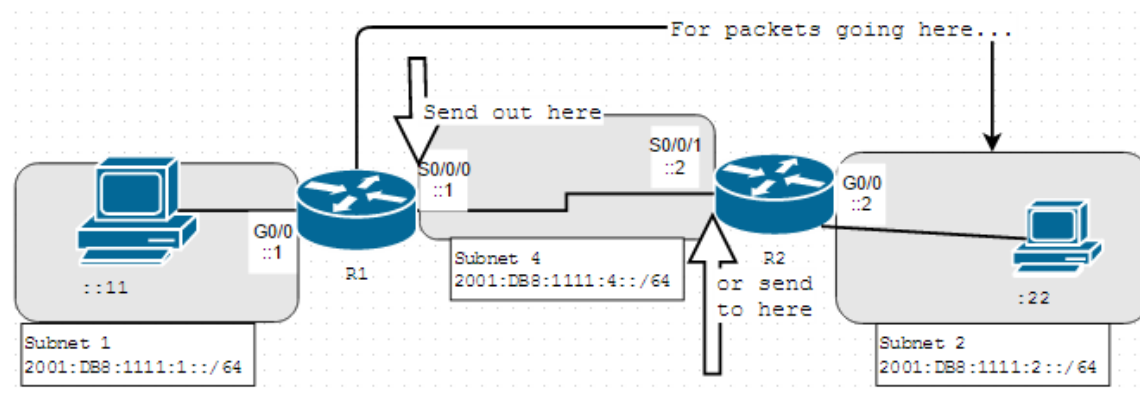
C 2001:db8:1111:5::/64 [0/0]

via Serial0/0/1, directly connected

show ipv6 route local

!Omitted

L 2001:db8:1111:1::1/128 [0/0]



ADD

!Static route on router R1

R1(config)# **ipv6 route 2001:db8:1111:2::/64 s0/0/0**

!Static route on router R2

R2(config)# **ipv6 route 2001:db8:1111:1::/64 s0/0/1**

VERIFY

R1(config)# **show ipv6 route static**

S 2001:db8:1111:2::/64 [1/0]

via Serial0/0/0, directly connected

R1(config)# **show ipv6 route 2001:db8:1111:2::22**

Routing entry for 2001:db8:1111:2::/64

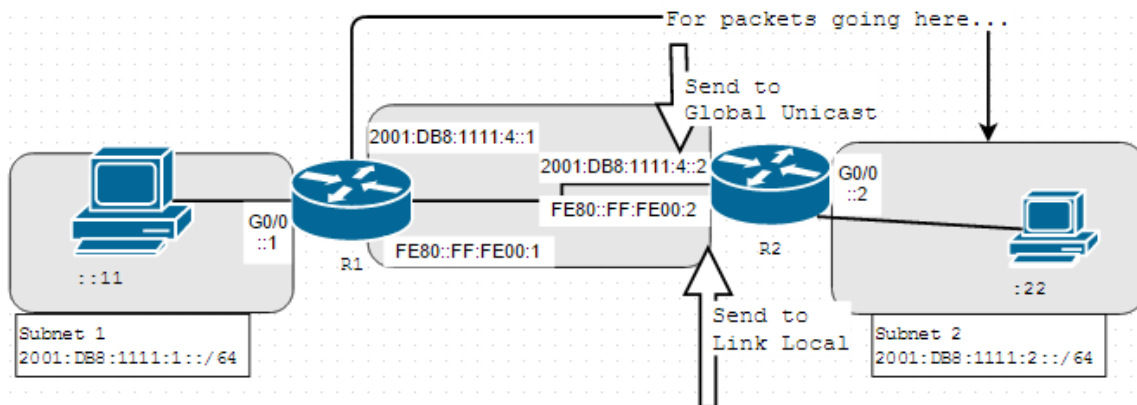
Known via "static", distance 1, metric 0

Route count is 1/1, share count 0

Routing paths:

directly connected via Serial0/0/0

Last updated 00:01:29 ago



Static Routes Using Next-Hop IPv6 Address: Using Unicast or Link-Local as Next-Hop Address

!First command is on R1, listing R2's global unicast

R1(config)#**ipv6 route 2001:db8:1111:2::/64 2001:db8:1111:4::2**

!This command is on R2, listing R1's global unicast

R2(config)#**ipv6 route 2001:db8:1111:1::/64 2001:db8:1111:4::1**

!Verify routes with **show ipv6 route static**

!First command is on R1, listing R2's link local address

R1(config)#**ipv6 route 2001:db8:1111:2::/64 s0/0/0 FE80:FF:FE00:2**

!This command is on R2, listing R1's link local address

R2(config)#**ipv6 route 2001:db8:1111:1::/64 s0/0/1 FE80:FF:FE00:1**

!Verify routes with **show ipv6 route static**

- **Static Default Routes**

- Tells the router what to do when a packet does not match a route
 - Without a default route, router discards packet
 - With a default route, router forwards packet to the default route

!Forward out B1's S0/0/1 local interface...

B1(config)#**ipv6 route ::0 S0/0/1**

!Static Default Route

B1#**show ipv6 route static**

!Omitted

S ::0 [1/0]

- Migration from OSPFv2 to OSPFv3
 - Before IPv6, the company supports IPv4 using OSPFv2
 - Company plans to use a dual-stack approach
 - Companies add OSPFv3 configuration to all the routers
 - OSPFv2 support stays in place

OSPFv2 Indirectly Enables OSPF on the Interface and OSPFv3 Configuration Directly Enables OSPF

```
router ospf 1
  router-id 1.1.1.1
  network 10.0.0.0 0.255.255.255 area 0 <--- indirectly enable on int
interface s0/0/0
  ip address 10.1.1.1 255.255.255.0
```

```
ipv6 router ospf 1
  router-id 1.1.1.1
```

```
interface s0/0/0
  ipv6 ospf 1 area 0    <----- directly enable on int
```

- **Configuring Single-Area OSPFv3**

- OSPFv2:
 - If **router-id rid** OSPF subcommand is configured, use it
 - If the RID is not set, check the loopback interfaces with up status
 - Choose the highest numeric IP address
 - If neither are used, router picks highest numeric IPv4 address
- OSPFv3:
 - Use **ipv6 router ospf process-id**
 - Ensure the router has an OSPF router ID
 - Configuring the **router-id id-value** router subcommand
 - Preconfiguring an IPv4 on any loopback who is up
 - Preconfiguring an IPv4 on any working interface who is up
 - Configure the **ipv6 ospf process-id area area-number** on each OSPF interface

ipv6 unicast-routing

interface serial0/0/0

no ip address

ipv6 address 2001:db8:1111:4::1/64

!

interface s0/0/1

no ip address

ipv6 address 2001:db8:1111:5::1/64

!

interface GigabitEthernet0/0

no ip address

ipv6 address 2001:db8:1111:5::1/64

!Enabling OSPFv3 on three interfaces

config t

ipv6 router ospf 1

router-id 1.1.1.1

interface s0/0/0

ipv6 ospf 1 area 0

interface s0/0/1

ipv6 ospf 1 area 0

int gi0/0

ipv6 ospf 1 area 0

end

ipv6 unicast-routing

ipv6 router ospf 2

router-id 2.2.2.2

!

int s0/0/1

ipv6 address 2001:db8:1111:4::2

ipv6 ospf 2 area 0

!

int gi0/0

ipv6 address 2001:db8:1111:2::2

ipv6 ospf 2 area 0

<i>To Display Details About...</i>	<i>OSPFv2</i>	<i>OSPFv3</i>
OSPF process	show ip ospf	show ipv6 ospf
All sources of routing info	show ip protocols	show ipv6 protocols
Details about OSPF-enabled interfaces	show ip ospf interface	show ipv6 ospf interface
Concise info on OSPF-enabled interfaces	show ip ospf interface brief	show ipv6 ospf interface brief
List of neighbors	show ip ospf neighbor	show ipv6 ospf neighbor
Summary of LSDB	show ip ospf database	show ipv6 ospf database
OSPF-learned routes	show ip route ospf	show ipv6 route ospf

R1#**show ipv6 ospf**

Routing Process "ospfv3 1" with ID 1.1.1.1

!Omitted

Reference bandwidth unit is 100 mbps

Area BACKBONE 0

Number of interfaces in this area is 3

!Verifying interfaces

R1# **show ipv6 ospf interface brief**

Interface	PID	Area	Intf ID	Cost	State	Nbrs	F/C
Gi0/0	1	0	3	1	DR	0/0	
Se0/0/1	1	0	7	64	P2P	1/1	
Se0/0/0	1	0	6	64	P2P	1/1	

R1# **show ipv6 protocols**

IPv6 Routing Protocol is "connected"

IPv6 Routing Protocol is "ND"

IPv6 Routing Protocol is "ospf 1"

Interfaces (Area 0):

GigabitEthernet0/0

Serial0/0/1

Serial0/0/0

!First command is from R1, listing R2 and R2

show ipv6 ospf neighbor

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
3.3.3.3	0	FULL/	- 00:00:39	6	Serial0/0/1
2.2.2.2	0	FULL/	- 00:00:31	7	Serial0/0/0

!This command is from R2, listing R1 and R3

show ipv6 ospf neighbor

Neighbor ID	Pri	State	Dead Time	Interface ID	Interface
1.1.1.1	0	FULL/	- 00:00:39	6	Serial0/0/1
3.3.3.3	0	FULL/	- 00:00:31	7	Gi0/0

(Verifying OSPFv3 LSDB on R1)

show ipv6 ospf database

OSPFv3 Router with ID (2.2.2.2) (Process ID 2)

Router Link States (Area 0)

ADV Router	Age	Seq#	Fragment ID	Link count	Bits
1.1.1.1	452	0x80000002	0	2	None

2.2.2.2	456	0x80000004	0	2	None
3.3.3.3	457	0x80000005	0	2	None

(from R2)

show ipv6 route ospf

IPv6 Routing Table - default - 9 entries

!Omitted

O 2001:DB8:1111:1::/64 [110/65]
via FE80::FF:FE00:1, serial0/0/1

ipv6 unicast routing -- for static routing

:4::1/64

:5::1/64

:5::1/64

2001:db8:1111

show ipv6 route [local | static |]

show ipv6 route (address) to show info specific to that route

ipv6 route (address) s0/0 (next hop) --add to table - can be link local or global unicast of destination port

static default routes

ipv6 route ::0 s0/0/1

S ::0 will show up in routing table

(::0 any address)

same: router ospf 1, router-id 1.1.1.1 - you should consider rid to be required in ipv6

OSPF instead of indirectly enabling on interfaces ipv6 has to have "ipv6 ospf 1 area 0"

(no "network ip mask area command for router part)

ch 27-29

Address	Description	Available Scopes
ff0X::1	All nodes address, identify the group of all IPv6 nodes	Available in scope 1 (interface-local) and 2 (link-local):

ff01::1 → All nodes in the interface-local

ff02::1 → All nodes in the link-local

ff0X::2 All routers Available in scope 1 (interface-local), 2 (link-local) and 5 (site-local):

ff01::2 → All routers in the interface-local

ff02::2 → All routers in the link-local

ff05::2 → All routers in the site-local

ff02::5	OSPF/IGP	2 (link-local)
ff02::6	OSPF/IGP Designated Routers	2 (link-local)
ff02::9	RIP Routers	2 (link-local)
ff02::a	EIGRP Routers	2 (link-local)

ff02::d	All PIM Routers	2 (link-local)
ff02::1a	All RPL Routers	2 (link-local)
ff0X::fb	mDNSv6	Available in all scopes
ff0X::101	All Network Time Protocol (NTP) servers	Available in all scopes
ff02::1:1	Link Name	2 (link-local)
ff02::1:2	All-dhcp-agents	2 (link-local)
ff02::1:3	Link-local Multicast Name Resolution	2 (link-local)
ff05::1:3	All-dhcp-servers	5 (site-local)
ff02::1:ff00:0/104	Solicited-node multicast address. See below	2 (link-local)
ff02::2:ff00:0/104	Node Information Queries	2 (link-local)