

Kubernetes - Kubectl: The Missing Manpages

The kubectl CLI is written in Golang, and uses Cobra to automatically generate documentation. You use 'kubectl help' and 'kubectl [COMMAND] --help' to get manpage-like output. Since these aren't available as manpages, they won't be found in various WWW manpage collections to read (so you can look at them on a device without Kubernetes installed like a phone). Not included here is output of 'kubectl explain' for anything, so it's just 'the missing manpages'

=====

```
kubectl help
kubectl controls the Kubernetes cluster manager.
```

Find more information at: <https://kubernetes.io/docs/reference/kubectl/>

Basic Commands (Beginner):

```
create          Create a resource from a file or from stdin
expose          Take a replication controller, service, deployment or pod and expose it as
a new Kubernetes service
run             Run a particular image on the cluster
set             Set specific features on objects
```

Basic Commands (Intermediate):

```
explain         Get documentation for a resource
get             Display one or many resources
edit            Edit a resource on the server
delete          Delete resources by file names, stdin, resources and names, or by
resources and label selector
```

Deploy Commands:

```
rollout         Manage the rollout of a resource
scale           Set a new size for a deployment, replica set, or replication controller
autoscale       Auto-scale a deployment, replica set, stateful set, or replication
controller
```

Cluster Management Commands:

```
certificate      Modify certificate resources.
cluster-info     Display cluster information
top              Display resource (CPU/memory) usage
cordon           Mark node as unschedulable
uncordon         Mark node as schedulable
drain            Drain node in preparation for maintenance
taint            Update the taints on one or more nodes
```

Troubleshooting and Debugging Commands:

```
describe        Show details of a specific resource or group of resources
logs            Print the logs for a container in a pod
attach          Attach to a running container
exec            Execute a command in a container
port-forward     Forward one or more local ports to a pod
proxy           Run a proxy to the Kubernetes API server
cp              Copy files and directories to and from containers
auth            Inspect authorization
debug           Create debugging sessions for troubleshooting workloads and nodes
events          List events
```

Advanced Commands:

```
diff            Diff the live version against a would-be applied version
apply           Apply a configuration to a resource by file name or stdin
patch           Update fields of a resource
replace         Replace a resource by file name or stdin
wait           Experimental: Wait for a specific condition on one or many resources
kustomize       Build a kustomization target from a directory or URL.
```

Settings Commands:

```
label           Update the labels on a resource
annotate        Update the annotations on a resource
```

completion Output shell completion code for the specified shell (bash, zsh, fish, or powershell)

Other Commands:

alpha Commands for features in alpha
api-resources Print the supported API resources on the server
api-versions Print the supported API versions on the server, in the form of "group/version"
config Modify kubeconfig files
plugin Provides utilities for interacting with plugins
version Print the client and server version information

Usage:

kubectl [flags] [options]

=====

[tm@freeipa ~]\$ kubectl options

The following options can be passed to any command:

--as='':

Username to impersonate for the operation. User could be a regular user or a service account in a namespace.

--as-group=[]:

Group to impersonate for the operation, this flag can be repeated to specify multiple groups.

--as-uid='':

UID to impersonate for the operation.

--cache-dir='/home/tm/.kube/cache':

Default cache directory

--certificate-authority='':

Path to a cert file for the certificate authority

--client-certificate='':

Path to a client certificate file for TLS

--client-key='':

Path to a client key file for TLS

--cluster='':

The name of the kubeconfig cluster to use

--context='':

The name of the kubeconfig context to use

--disable-compression=false:

If true, opt-out of response compression for all requests to the server

--insecure-skip-tls-verify=false:

If true, the server's certificate will not be checked for validity. This will make your HTTPS connections insecure

--kubeconfig='':

Path to the kubeconfig file to use for CLI requests.

--log-flush-frequency=5s:

Maximum number of seconds between log flushes

--match-server-version=false:

Require server version to match client version

-n, --namespace='':
 If present, the namespace scope for this CLI request

--password='':
 Password for basic authentication to the API server

--profile='none':
 Name of profile to capture. One of (none|cpu|heap|goroutine|threadcreate|block|mutex)

--profile-output='profile.pprof':
 Name of the file to write the profile to

--request-timeout='0':
 The length of time to wait before giving up on a single server request. Non-zero values should contain a corresponding time unit (e.g. 1s, 2m, 3h). A value of zero means don't timeout requests.

-s, --server='':
 The address and port of the Kubernetes API server

--tls-server-name='':
 Server name to use for server certificate validation. If it is not provided, the hostname used to contact the server is used

--token='':
 Bearer token for authentication to the API server

--user='':
 The name of the kubeconfig user to use

--username='':
 Username for basic authentication to the API server

-v, --v=0:
 number for the log level verbosity

--vmodule=:
 comma-separated list of pattern=N settings for file-filtered logging (only works for the default text log format)

--warnings-as-errors=false:
 Treat warnings received from the server as errors and exit with a non-zero exit code

=====

kubectl create --help
 Create a resource from a file or from stdin.

JSON and YAML formats are accepted.

Examples:

Create a pod using the data in pod.json
 kubectl create -f ./pod.json

Create a pod based on the JSON passed into stdin
 cat pod.json | kubectl create -f -

```
# Edit the data in registry.yaml in JSON then create the resource using the edited data
kubectl create -f registry.yaml --edit -o json
```

Available Commands:

clusterrole	Create a cluster role
clusterrolebinding	Create a cluster role binding for a particular cluster role
configmap	Create a config map from a local file, directory or literal value
cronjob	Create a cron job with the specified name
deployment	Create a deployment with the specified name
ingress	Create an ingress with the specified name
job	Create a job with the specified name
namespace	Create a namespace with the specified name
poddisruptionbudget	Create a pod disruption budget with the specified name
priorityclass	Create a priority class with the specified name
quota	Create a quota with the specified name
role	Create a role with single rule
rolebinding	Create a role binding for a particular role or cluster role
secret	Create a secret using specified subcommand
service	Create a service using a specified subcommand
serviceaccount	Create a service account with the specified name
token	Request a service account token

Options:

```
--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to goyaml and jsonpath output
    formats.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

--edit=false:
    Edit the API resource before creating

--field-manager='kubectl-create':
    Name of the manager used to track field ownership.

-f, --filename=[ ]:
    Filename, directory, or URL to files to use to create the resource

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
    with -f or -R.

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file).

--raw='':
    Raw URI to POST to the server. Uses the transport specified by the
    kubeconfig file.

-R, --recursive=false:
    Process the directory used in -f, --filename recursively. Useful when
    you want to manage related manifests organized within the same
    directory.

--save-config=false:
    If true, the configuration of current object will be saved in its
    annotation. Otherwise, the annotation will be unchanged. This flag is
    useful when you want to perform kubectl apply on this object in the
    future.
```

```

-l, --selector='':
    Selector (label query) to filter on, supports '=', '==', and
    '!='.(e.g. -l key1=value1,key2=value2). Matching objects must satisfy
    all of the specified label constraints.

--show-managed-fields=false:
    If true, keep the managedFields when printing objects in JSON or YAML
    format.

--template='':
    Template string or path to template file to use when -o=go-template,
    -o=go-template-file. The template format is go lang templates
    [http://golang.org/pkg/text/template/#pkg-overview].

--validate='strict':
    Must be one of: strict (or true), warn, ignore (or false).
        "true" or "strict" will use a schema to validate the input and fail the request
        if invalid. It will perform server side validation if
        ServerSideFieldValidation is enabled on the api-server, but will fall
        back to less reliable client-side validation if not.
        "warn" will
        warn about unknown or duplicate fields without blocking the request if
        server-side field validation is enabled on the API server, and behave
        as "ignore" otherwise.
        "false" or "ignore" will not perform any
        schema validation, silently dropping any unknown or duplicate fields.

--windows-line-endings=false:
    Only relevant if --edit=true. Defaults to the line ending native to
    your platform.

```

Usage:

```
kubectl create -f FILENAME [options]
```

=====

```
kubectl expose --help
```

Expose a resource as a new Kubernetes service.

Looks up a deployment, service, replica set, replication controller or pod by name and uses the selector for that resource as the selector for a new service on the specified port. A deployment or replica set will be exposed as a service only if its selector is convertible to a selector that service supports, i.e. when the selector contains only the matchLabels component. Note that if no port is specified via --port and the exposed resource has multiple ports, all will be re-used by the new service. Also if no labels are specified, the new service will re-use the labels from the resource it exposes.

Possible resources include (case insensitive):

```
pod (po), service (svc), replicationcontroller (rc), deployment (deploy),
replicaset (rs)
```

Examples:

```
# Create a service for a replicated nginx, which serves on port 80 and connects to the
containers on port 8000
```

```
kubectl expose rc nginx --port=80 --target-port=8000
```

```
# Create a service for a replication controller identified by type and name specified in
"nginx-controller.yaml", which serves on port 80 and connects to the containers on port 8000
```

```
kubectl expose -f nginx-controller.yaml --port=80 --target-port=8000
```

```
# Create a service for a pod valid-pod, which serves on port 444 with the name "frontend"
```

```

kubect1 expose pod valid-pod --port=444 --name=frontend

# Create a second service based on the above service, exposing the container port 8443 as
port 443 with the name "nginx-https"
kubect1 expose service nginx --port=443 --target-port=8443 --name=nginx-https

# Create a service for a replicated streaming application on port 4100 balancing UDP
traffic and named 'video-stream'.
kubect1 expose rc streamer --port=4100 --protocol=UDP --name=video-stream

# Create a service for a replicated nginx using replica set, which serves on port 80 and
connects to the containers on port 8000
kubect1 expose rs nginx --port=80 --target-port=8000

# Create a service for an nginx deployment, which serves on port 80 and connects to the
containers on port 8000
kubect1 expose deployment nginx --port=80 --target-port=8000

```

Options:

```

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go1.10 and jsonpath output
    formats.

--cluster-ip='':
    ClusterIP to be assigned to the service. Leave empty to auto-allocate,
    or set to 'None' to create a headless service.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

--external-ip='':
    Additional external IP address (not managed by Kubernetes) to accept
    for the service. If this IP is routed to a node, the service can be
    accessed by this IP in addition to its generated service IP.

--field-manager='kubect1-expose':
    Name of the manager used to track field ownership.

-f, --filename=[ ]:
    Filename, directory, or URL to files identifying the resource to
    expose a service

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
    with -f or -R.

-l, --labels='':
    Labels to apply to the service created by this call.

--load-balancer-ip='':
    IP to assign to the LoadBalancer. If empty, an ephemeral IP will be
    created and used (cloud-provider specific).

--name='':
    The name for the newly created object.

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file).

--override-type='merge':
    The method used to override the generated object: json, merge, or

```

```

    strategic.

--overrides='':
    An inline JSON override for the generated object. If this is
    non-empty, it is used to override the generated object. Requires that
    the object supply a valid apiVersion field.

--port='':
    The port that the service should serve on. Copied from the resource
    being exposed, if unspecified

--protocol='':
    The network protocol for the service to be created. Default is 'TCP'.

-R, --recursive=false:
    Process the directory used in -f, --filename recursively. Useful when
    you want to manage related manifests organized within the same
    directory.

--save-config=false:
    If true, the configuration of current object will be saved in its
    annotation. Otherwise, the annotation will be unchanged. This flag is
    useful when you want to perform kubectl apply on this object in the
    future.

--selector='':
    A label selector to use for this service. Only equality-based selector
    requirements are supported. If empty (the default) infer the selector
    from the replication controller or replica set.)

--session-affinity='':
    If non-empty, set the session affinity for the service to this; legal
    values: 'None', 'ClientIP'

--show-managed-fields=false:
    If true, keep the managedFields when printing objects in JSON or YAML
    format.

--target-port='':
    Name or number for the port on the container that the service should
    direct traffic to. Optional.

--template='':
    Template string or path to template file to use when -o=go-template,
    -o=go-template-file. The template format is go templates
    [http://golang.org/pkg/text/template/#pkg-overview].

--type='':
    Type for this service: ClusterIP, NodePort, LoadBalancer, or
    ExternalName. Default is 'ClusterIP'.

```

Usage:

```

kubectl expose (-f FILENAME | TYPE NAME) [--port=port] [--protocol=TCP|UDP|SCTP] [--
target-port=number-or-name] [--name=name] [--external-ip=external-ip-of-service] [--
type=type] [options]

```

```

=====
kubectl run --help

```

Create and run a particular image in a pod.

Examples:

```

# Start a nginx pod
kubectl run nginx --image=nginx

```

```

# Start a hazelcast pod and let the container expose port 5701

```

```

kubect1 run hazelcast --image=hazelcast/hazelcast --port=5701

# Start a hazelcast pod and set environment variables "DNS_DOMAIN=cluster" and
"POD_NAMESPACE=default" in the container
kubect1 run hazelcast --image=hazelcast/hazelcast --env="DNS_DOMAIN=cluster" --
env="POD_NAMESPACE=default"

# Start a hazelcast pod and set labels "app=hazelcast" and "env=prod" in the container
kubect1 run hazelcast --image=hazelcast/hazelcast --labels="app=hazelcast,env=prod"

# Dry run; print the corresponding API objects without creating them
kubect1 run nginx --image=nginx --dry-run=client

# Start a nginx pod, but overload the spec with a partial set of values parsed from JSON
kubect1 run nginx --image=nginx --overrides='{ "apiVersion": "v1", "spec": { ... } }'

# Start a busybox pod and keep it in the foreground, don't restart it if it exits
kubect1 run -i -t busybox --image=busybox --restart=Never

# Start the nginx pod using the default command, but use custom arguments (arg1 .. argN)
for that command
kubect1 run nginx --image=nginx -- <arg1> <arg2> ... <argN>

# Start the nginx pod using a different command and custom arguments
kubect1 run nginx --image=nginx --command -- <cmd> <arg1> ... <argN>

```

Options:

```

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go lang and jsonpath output
    formats.

--annotations=[ ]:
    Annotations to apply to the pod.

--attach=false:
    If true, wait for the Pod to start running, and then attach to the Pod
    as if 'kubect1 attach ...' were called. Default false, unless
    '-i/--stdin' is set, in which case the default is true. With
    '--restart=Never' the exit code of the container process is returned.

--command=false:
    If true and extra arguments are present, use them as the 'command'
    field in the container, rather than the 'args' field which is the
    default.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

--env=[ ]:
    Environment variables to set in the container.

--expose=false:
    If true, create a ClusterIP service associated with the pod. Requires
    '--port'.

--field-manager='kubect1-run':
    Name of the manager used to track field ownership.

--image='':
    The image for the container to run.

--image-pull-policy='':
    The image pull policy for the container. If left empty, this value

```


will not be specified by the client and defaulted by the server.

-l, --labels='':
Comma separated labels to apply to the pod. Will override previous values.

--leave-stdin-open=false:
If the pod is started in interactive mode or with stdin, leave stdin open after the first attach completes. By default, stdin will be closed after the first attach completes.

-o, --output='':
Output format. One of: (json, yaml, name, go-template, go-template-file, template, templatefile, jsonpath, jsonpath-as-json, jsonpath-file).

--override-type='merge':
The method used to override the generated object: json, merge, or strategic.

--overrides='':
An inline JSON override for the generated object. If this is non-empty, it is used to override the generated object. Requires that the object supply a valid apiVersion field.

--pod-running-timeout=1m0s:
The length of time (like 5s, 2m, or 3h, higher than zero) to wait until at least one pod is running

--port='':
The port that this container exposes.

--privileged=false:
If true, run the container in privileged mode.

-q, --quiet=false:
If true, suppress prompt messages.

--restart='Always':
The restart policy for this Pod. Legal values [Always, OnFailure, Never].

--rm=false:
If true, delete the pod after it exits. Only valid when attaching to the container, e.g. with '--attach' or with '-i/--stdin'.

--save-config=false:
If true, the configuration of current object will be saved in its annotation. Otherwise, the annotation will be unchanged. This flag is useful when you want to perform kubectl apply on this object in the future.

--show-managed-fields=false:
If true, keep the managedFields when printing objects in JSON or YAML format.

-i, --stdin=false:
Keep stdin open on the container in the pod, even if nothing is attached.

--template='':
Template string or path to template file to use when -o=go-template, -o=go-template-file. The template format is golang templates [<http://golang.org/pkg/text/template/#pkg-overview>].

-t, --tty=false:

Allocate a TTY for the container in the pod.

Usage:

```
kubectl run NAME --image=image [--env="key=value"] [--port=port] [--dry-run=server|client]
[--overrides=inline-json] [--command] -- [COMMAND] [args...] [options]
```

```
=====
kubectl set --help
```

Configure application resources.

These commands help you make changes to existing application resources.

Available Commands:

env	Update environment variables on a pod template
image	Update the image of a pod template
resources	Update resource requests/limits on objects with pod templates
selector	Set the selector on a resource
serviceaccount	Update the service account of a resource
subject	Update the user, group, or service account in a role binding or cluster role binding

Usage:

```
kubectl set SUBCOMMAND [options]
```

```
=====
kubectl get --help
```

Display one or many resources.

Prints a table of the most important information about the specified resources. You can filter the list using a label selector and the `--selector` flag. If the desired resource type is namespaced you will only see results in your current namespace unless you pass `--all-namespaces`.

By specifying the output as 'template' and providing a Go template as the value of the `--template` flag, you can filter the attributes of the fetched resources.

Use "kubectl api-resources" for a complete list of supported resources.

Examples:

```
# List all pods in ps output format
kubectl get pods
```

```
# List all pods in ps output format with more information (such as node name)
kubectl get pods -o wide
```

```
# List a single replication controller with specified NAME in ps output format
kubectl get replicationcontroller web
```

```
# List deployments in JSON output format, in the "v1" version of the "apps" API group
kubectl get deployments.v1.apps -o json
```

```
# List a single pod in JSON output format
kubectl get -o json pod web-pod-13je7
```

```
# List a pod identified by type and name specified in "pod.yaml" in JSON output format
kubectl get -f pod.yaml -o json
```

```
# List resources from a directory with kustomization.yaml - e.g. dir/kustomization.yaml
kubectl get -k dir/
```

```
# Return only the phase value of the specified pod
kubectl get -o template pod/web-pod-13je7 --template={{.status.phase}}
```

```
# List resource information in custom columns
```

```

kubect1 get pod test-pod -o custom-
columns=CONTAINER:.spec.containers[0].name,IMAGE:.spec.containers[0].image

# List all replication controllers and services together in ps output format
kubect1 get rc,services

# List one or more resources by their type and names
kubect1 get rc/web service/frontend pods/web-pod-13je7

# List status subresource for a single pod.
kubect1 get pod web-pod-13je7 --subresource status

```

Options:

```

-A, --all-namespaces=false:
    If present, list the requested object(s) across all namespaces.
    Namespace in current context is ignored even if specified with
    --namespace.

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go1ang and jsonpath output
    formats.

--chunk-size=500:
    Return large lists in chunks rather than all at once. Pass 0 to
    disable. This flag is beta and may change in the future.

--field-selector='':
    Selector (field query) to filter on, supports '=', '==', and
    '!='. (e.g. --field-selector key1=value1,key2=value2). The server only
    supports a limited number of field queries per type.

-f, --filename=[ ]:
    Filename, directory, or URL to files identifying the resource to get
    from a server.

--ignore-not-found=false:
    If the requested object does not exist the command will return exit
    code 0.

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
    with -f or -R.

-L, --label-columns=[ ]:
    Accepts a comma separated list of labels that are going to be
    presented as columns. Names are case-sensitive. You can also use
    multiple flag options like -L label1 -L label2...

--no-headers=false:
    When using the default or custom-column output format, don't print
    headers (default print headers).

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file, custom-columns, custom-columns-file, wide). See custom
    columns
    [https://kubernetes.io/docs/reference/kubect1/#custom-columns], go1ang
    template [http://go1ang.org/pkg/text/template/#pkg-overview] and
    jsonpath template
    [https://kubernetes.io/docs/reference/kubect1/jsonpath/].

--output-watch-events=false:
    Output watch event objects when --watch or --watch-only is used.
    Existing objects are output as initial ADDED events.

```

--raw='':
Raw URI to request from the server. Uses the transport specified by the kubeconfig file.

-R, --recursive=false:
Process the directory used in -f, --filename recursively. Useful when you want to manage related manifests organized within the same directory.

-l, --selector='':
Selector (label query) to filter on, supports '=', '==', and '!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy all of the specified label constraints.

--server-print=true:
If true, have the server return the appropriate table output. Supports extension APIs and CRDs.

--show-kind=false:
If present, list the resource type for the requested object(s).

--show-labels=false:
When printing, show all labels as the last column (default hide labels column)

--show-managed-fields=false:
If true, keep the managedFields when printing objects in JSON or YAML format.

--sort-by='':
If non-empty, sort list types using this field specification. The field specification is expressed as a JSONPath expression (e.g. '{.metadata.name}'). The field in the API resource specified by this JSONPath expression must be an integer or a string.

--subresource='':
If specified, gets the subresource of the requested object. Must be one of [status scale]. This flag is alpha and may change in the future.

--template='':
Template string or path to template file to use when -o=go-template, -o=go-template-file. The template format is go lang templates [http://golang.org/pkg/text/template/#pkg-overview].

-w, --watch=false:
After listing/getting the requested object, watch for changes.

--watch-only=false:
Watch for changes to the requested object(s), without listing/getting first.

Usage:

```
kubectl get
[(-o|--output=)json|yaml|name|go-template|go-template-file|template|templatefile|jsonpath|
jsonpath-as-json|jsonpath-file|custom-columns|custom-columns-file|wide] (TYPE[.VERSION]
[.GROUP] [NAME | -l label] | TYPE[.VERSION][.GROUP]/NAME ...) [flags] [options]
```

=====

kubectl explain --help
List the fields for supported resources.

This command describes the fields associated with each supported API resource.

Fields are identified via a simple JSONPath identifier:

```
<type>.<fieldName>[.<fieldName>]
```

Add the `--recursive` flag to display all of the fields at once without descriptions. Information about each field is retrieved from the server in OpenAPI format.

Use `"kubectl api-resources"` for a complete list of supported resources.

Examples:

```
# Get the documentation of the resource and its fields
kubectl explain pods

# Get the documentation of a specific field of a resource
kubectl explain pods.spec.containers
```

Options:

```
--api-version='':
    Get different explanations for particular API version (API
    group/version)

--recursive=false:
    Print the fields of fields (Currently only 1 level deep)
```

Usage:

```
kubectl explain RESOURCE [options]
```

```
=====

kubectl edit --help
```

Edit a resource from the default editor.

The `edit` command allows you to directly edit any API resource you can retrieve via the command-line tools. It will open the editor defined by your `KUBE_EDITOR`, or `EDITOR` environment variables, or fall back to `'vi'` for Linux or `'notepad'` for Windows. You can edit multiple objects, although changes are applied one at a time. The command accepts file names as well as command-line arguments, although the files you point to must be previously saved versions of resources.

Editing is done with the API version used to fetch the resource. To edit using a specific API version, fully-qualify the resource, version, and group.

The default format is YAML. To edit in JSON, specify `"-o json"`.

The flag `--windows-line-endings` can be used to force Windows line endings, otherwise the default for your operating system will be used.

In the event an error occurs while updating, a temporary file will be created on disk that contains your unapplied changes. The most common error when updating a resource is another editor changing the resource on the server. When this occurs, you will have to apply your changes to the newer version of the resource, or update your temporary saved copy to include the latest resource version.

Examples:

```
# Edit the service named 'registry'
kubectl edit svc/registry

# Use an alternative editor
KUBE_EDITOR="nano" kubectl edit svc/registry

# Edit the job 'myjob' in JSON using the v1 API format
kubectl edit job.v1.batch/myjob -o json
```

```
# Edit the deployment 'mydeployment' in YAML and save the modified config in its
annotation
kubectl edit deployment/mydeployment -o yaml --save-config

# Edit the deployment/mydeployment's status subresource
kubectl edit deployment mydeployment --subresource='status'
```

Options:

```
--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go lang and jsonpath output
    formats.

--field-manager='kubectl-edit':
    Name of the manager used to track field ownership.

-f, --filename=[ ]:
    Filename, directory, or URL to files to use to edit the resource

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
    with -f or -R.

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file).

--output-patch=false:
    Output the patch if the resource is edited.

-R, --recursive=false:
    Process the directory used in -f, --filename recursively. Useful when
    you want to manage related manifests organized within the same
    directory.

--save-config=false:
    If true, the configuration of current object will be saved in its
    annotation. Otherwise, the annotation will be unchanged. This flag is
    useful when you want to perform kubectl apply on this object in the
    future.

--show-managed-fields=false:
    If true, keep the managedFields when printing objects in JSON or YAML
    format.

--subresource='':
    If specified, edit will operate on the subresource of the requested
    object. Must be one of [status]. This flag is alpha and may change in
    the future.

--template='':
    Template string or path to template file to use when -o=go-template,
    -o=go-template-file. The template format is go lang templates
    [http://golang.org/pkg/text/template/#pkg-overview].

--validate='strict':
    Must be one of: strict (or true), warn, ignore (or false).
    "true" or
    "strict" will use a schema to validate the input and fail the request
    if invalid. It will perform server side validation if
    ServerSideFieldValidation is enabled on the api-server, but will fall
    back to less reliable client-side validation if not.
    "warn" will
    warn about unknown or duplicate fields without blocking the request if
```

server-side field validation is enabled on the API server, and behave as "ignore" otherwise.

"false" or "ignore" will not perform any schema validation, silently dropping any unknown or duplicate fields.

--windows-line-endings=false:

Defaults to the line ending native to your platform.

Usage:

kubect1 edit (RESOURCE/NAME | -f FILENAME) [options]

=====

kubect1 delete --help

Delete resources by file names, stdin, resources and names, or by resources and label selector.

JSON and YAML formats are accepted. Only one type of argument may be specified: file names, resources and names, or resources and label selector.

Some resources, such as pods, support graceful deletion. These resources define a default period before they are forcibly terminated (the grace period) but you may override that value with the --grace-period flag, or pass --now to set a grace-period of 1. Because these resources often represent entities in the cluster, deletion may not be acknowledged immediately. If the node hosting a pod is down or cannot reach the API server, termination may take significantly longer than the grace period. To force delete a resource, you must specify the --force flag. Note: only a subset of resources support graceful deletion. In absence of the support, the --grace-period flag is ignored.

IMPORTANT: Force deleting pods does not wait for confirmation that the pod's processes have been terminated, which can leave those processes running until the node detects the deletion and completes graceful deletion. If your processes use shared storage or talk to a remote API and depend on the name of the pod to identify themselves, force deleting those pods may result in multiple processes running on different machines using the same identification which may lead to data corruption or inconsistency. Only force delete pods when you are sure the pod is terminated, or if your application can tolerate multiple copies of the same pod running at once. Also, if you force delete pods, the scheduler may place new pods on those nodes before the node has released those resources and causing those pods to be evicted immediately.

Note that the delete command does NOT do resource version checks, so if someone submits an update to a resource right when you submit a delete, their update will be lost along with the rest of the resource.

After a CustomResourceDefinition is deleted, invalidation of discovery cache may take up to 6 hours. If you don't want to wait, you might want to run "kubect1 api-resources" to refresh the discovery cache.

Examples:

Delete a pod using the type and name specified in pod.json
kubect1 delete -f ./pod.json

Delete resources from a directory containing kustomization.yaml - e.g.
dir/kustomization.yaml
kubect1 delete -k dir

Delete resources from all files that end with '.json' - i.e. expand wildcard characters in file names
kubect1 delete -f '*.json'

Delete a pod based on the type and name in the JSON passed into stdin
cat pod.json | kubect1 delete -f -

```
# Delete pods and services with same names "baz" and "foo"
kubectl delete pod,service baz foo

# Delete pods and services with label name=myLabel
kubectl delete pods,services -l name=myLabel

# Delete a pod with minimal delay
kubectl delete pod foo --now

# Force delete a pod on a dead node
kubectl delete pod foo --force

# Delete all pods
kubectl delete pods --all
```

Options:

```
--all=false:
    Delete all resources, in the namespace of the specified resource
    types.

-A, --all-namespaces=false:
    If present, list the requested object(s) across all namespaces.
    Namespace in current context is ignored even if specified with
    --namespace.

--cascade='background':
    Must be "background", "orphan", or "foreground". Selects the deletion
    cascading strategy for the dependents (e.g. Pods created by a
    ReplicationController). Defaults to background.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

--field-selector='':
    Selector (field query) to filter on, supports '=', '==', and
    '!='. (e.g. --field-selector key1=value1,key2=value2). The server only
    supports a limited number of field queries per type.

-f, --filename=[ ]:
    containing the resource to delete.

--force=false:
    If true, immediately remove resources from API and bypass graceful
    deletion. Note that immediate deletion of some resources may result in
    inconsistency or data loss and requires confirmation.

--grace-period=-1:
    Period of time in seconds given to the resource to terminate
    gracefully. Ignored if negative. Set to 1 for immediate shutdown. Can
    only be set to 0 when --force is true (force deletion).

--ignore-not-found=false:
    Treat "resource not found" as a successful delete. Defaults to "true"
    when --all is specified.

-k, --kustomize='':
    Process a kustomization directory. This flag can't be used together
    with -f or -R.

--now=false:
    If true, resources are signaled for immediate shutdown (same as
    --grace-period=1).

-o, --output='':
```


Output mode. Use "-o name" for shorter output (resource/name).

--raw='':

Raw URI to DELETE to the server. Uses the transport specified by the kubeconfig file.

-R, --recursive=false:

Process the directory used in -f, --filename recursively. Useful when you want to manage related manifests organized within the same directory.

-l, --selector='':

Selector (label query) to filter on, supports '=', '==', and '!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy all of the specified label constraints.

--timeout=0s:

The length of time to wait before giving up on a delete, zero means determine a timeout from the size of the object

--wait=true:

If true, wait for resources to be gone before returning. This waits for finalizers.

Usage:

kubectl delete ([-f FILENAME] | [-k DIRECTORY] | TYPE [(NAME | -l label | --all)])
[options]

=====

kubectl rollout --help

Manage the rollout of one or many resources.

Valid resource types include:

- * deployments
- * daemonsets
- * statefulsets

Examples:

Rollback to the previous deployment
kubectl rollout undo deployment/abc

Check the rollout status of a daemonset
kubectl rollout status daemonset/foo

Restart a deployment
kubectl rollout restart deployment/abc

Restart deployments with the app=nginx label
kubectl rollout restart deployment --selector=app=nginx

Available Commands:

history	View rollout history
pause	Mark the provided resource as paused
restart	Restart a resource
resume	Resume a paused resource
status	Show the status of the rollout
undo	Undo a previous rollout

Usage:

kubectl rollout SUBCOMMAND [options]

=====

kubectl scale --help
Set a new size for a deployment, replica set, replication controller, or stateful set.

Scale also allows users to specify one or more preconditions for the scale action.

If --current-replicas or --resource-version is specified, it is validated before the scale is attempted, and it is guaranteed that the precondition holds true when the scale is sent to the server.

Examples:

```
# Scale a replica set named 'foo' to 3
kubectl scale --replicas=3 rs/foo

# Scale a resource identified by type and name specified in "foo.yaml" to 3
kubectl scale --replicas=3 -f foo.yaml

# If the deployment named mysql's current size is 2, scale mysql to 3
kubectl scale --current-replicas=2 --replicas=3 deployment/mysql

# Scale multiple replication controllers
kubectl scale --replicas=5 rc/foo rc/bar rc/baz

# Scale stateful set named 'web' to 3
kubectl scale --replicas=3 statefulset/web
```

Options:

```
--all=false:
    Select all resources in the namespace of the specified resource types

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go lang and jsonpath output
    formats.

--current-replicas=-1:
    Precondition for current size. Requires that the current size of the
    resource match this value in order to scale. -1 (default) for no
    condition.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

-f, --filename=[ ]:
    Filename, directory, or URL to files identifying the resource to set a
    new size

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
    with -f or -R.

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file).

-R, --recursive=false:
    Process the directory used in -f, --filename recursively. Useful when
    you want to manage related manifests organized within the same
    directory.

--replicas=0:
    The new desired number of replicas. Required.
```

`--resource-version='':`
 Precondition for resource version. Requires that the current resource version match this value in order to scale.

`-l, --selector='':`
 Selector (label query) to filter on, supports '=', '==', and '!='. (e.g. `-l key1=value1,key2=value2`). Matching objects must satisfy all of the specified label constraints.

`--show-managed-fields=false:`
 If true, keep the managedFields when printing objects in JSON or YAML format.

`--template='':`
 Template string or path to template file to use when `-o=go-template`, `-o=go-template-file`. The template format is go lang templates [<http://golang.org/pkg/text/template/#pkg-overview>].

`--timeout=0s:`
 The length of time to wait before giving up on a scale operation, zero means don't wait. Any other values should contain a corresponding time unit (e.g. 1s, 2m, 3h).

Usage:

```
kubectl scale [--resource-version=version] [--current-replicas=count] --replicas=COUNT (-f
FILENAME | TYPE NAME) [options]
```

```
=====
kubectl autoscale --help
Creates an autoscaler that automatically chooses and sets the number of pods that run in a
Kubernetes cluster.
```

Looks up a deployment, replica set, stateful set, or replication controller by name and creates an autoscaler that uses the given resource as a reference. An autoscaler can automatically increase or decrease number of pods deployed within the system as needed.

Examples:

```
# Auto scale a deployment "foo", with the number of pods between 2 and 10, no target CPU
utilization specified so a default autoscaling policy will be used
kubectl autoscale deployment foo --min=2 --max=10

# Auto scale a replication controller "foo", with the number of pods between 1 and 5,
target CPU utilization at 80%
kubectl autoscale rc foo --max=5 --cpu-percent=80
```

Options:

`--allow-missing-template-keys=true:`
 If true, ignore any errors in templates when a field or map key is missing in the template. Only applies to go lang and jsonpath output formats.

`--cpu-percent=-1:`
 The target average CPU utilization (represented as a percent of requested CPU) over all the pods. If it's not specified or negative, a default autoscaling policy will be used.

`--dry-run='none':`
 Must be "none", "server", or "client". If client strategy, only print

the object that would be sent, without sending it. If server strategy, submit server-side request without persisting the resource.

`--field-manager='kubectl-autoscale':`
Name of the manager used to track field ownership.

`-f, --filename=[]:`
Filename, directory, or URL to files identifying the resource to autoscale.

`-k, --kustomize='':`
Process the kustomization directory. This flag can't be used together with `-f` or `-R`.

`--max=-1:`
The upper limit for the number of pods that can be set by the autoscaler. Required.

`--min=-1:`
The lower limit for the number of pods that can be set by the autoscaler. If it's not specified or negative, the server will apply a default value.

`--name='':`
The name for the newly created object. If not specified, the name of the input resource will be used.

`-o, --output='':`
Output format. One of: (json, yaml, name, go-template, go-template-file, template, templatefile, jsonpath, jsonpath-as-json, jsonpath-file).

`-R, --recursive=false:`
Process the directory used in `-f, --filename` recursively. Useful when you want to manage related manifests organized within the same directory.

`--save-config=false:`
If true, the configuration of current object will be saved in its annotation. Otherwise, the annotation will be unchanged. This flag is useful when you want to perform kubectl apply on this object in the future.

`--show-managed-fields=false:`
If true, keep the managedFields when printing objects in JSON or YAML format.

`--template='':`
Template string or path to template file to use when `-o=go-template, -o=go-template-file`. The template format is golang templates [<http://golang.org/pkg/text/template/#pkg-overview>].

Usage:

`kubectl autoscale (-f FILENAME | TYPE NAME | TYPE/NAME) [--min=MINPODS] --max=MAXPODS [--cpu-percent=CPU] [options]`

=====

`kubectl certificate --help`
Modify certificate resources.

Available Commands:

<code>approve</code>	Approve a certificate signing request
<code>deny</code>	Deny a certificate signing request

Usage:

```
kubectl certificate SUBCOMMAND [options]
```

```
=====

kubectl cluster-info --help
Display addresses of the control plane and services with label
kubernetes.io/cluster-service=true. To further debug and diagnose cluster
problems, use 'kubectl cluster-info dump'.
```

Examples:

```
# Print the address of the control plane and cluster services
kubectl cluster-info
```

Available Commands:

```
dump          Dump relevant information for debugging and diagnosis
```

Usage:

```
kubectl cluster-info [flags] [options]
```

```
=====

kubectl top --help
Display Resource (CPU/Memory) usage.
```

The top command allows you to see the resource consumption for nodes or pods.

This command requires Metrics Server to be correctly configured and working on the server.

Available Commands:

```
node          Display resource (CPU/memory) usage of nodes
pod           Display resource (CPU/memory) usage of pods
```

Usage:

```
kubectl top [flags] [options]
```

```
=====

kubectl cordon --help
Mark node as unschedulable.
```

Examples:

```
# Mark node "foo" as unschedulable
kubectl cordon foo
```

Options:

```
--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

-l, --selector='':
    Selector (label query) to filter on, supports '=', '==', and
    '!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy
    all of the specified label constraints.
```

Usage:

```
kubectl cordon NODE [options]
```

```
=====

kubectl uncordon --help
Mark node as schedulable.
```

Examples:

```
# Mark node "foo" as schedulable
kubectl uncordon foo
```

Options:

```
--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

-l, --selector='':
    Selector (label query) to filter on, supports '=', '==', and
    '!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy
    all of the specified label constraints.
```

Usage:

```
kubectl uncordon NODE [options]
```

=====

```
kubectl drain --help
```

Drain node in preparation for maintenance.

The given node will be marked unschedulable to prevent new pods from arriving. 'drain' evicts the pods if the API server supports <https://kubernetes.io/docs/concepts/workloads/pods/disruptions/> eviction. Otherwise, it will use normal DELETE to delete the pods. The 'drain' evicts or deletes all pods except mirror pods (which cannot be deleted through the API server). If there are daemon set-managed pods, drain will not proceed without `--ignore-daemonsets`, and regardless it will not delete any daemon set-managed pods, because those pods would be immediately replaced by the daemon set controller, which ignores unschedulable markings. If there are any pods that are neither mirror pods nor managed by a replication controller, replica set, daemon set, stateful set, or job, then drain will not delete any pods unless you use `--force`. `--force` will also allow deletion to proceed if the managing resource of one or more pods is missing.

'drain' waits for graceful termination. You should not operate on the machine until the command completes.

When you are ready to put the node back into service, use `kubectl uncordon`, which will make the node schedulable again.

https://kubernetes.io/images/docs/kubectl_drain.svg
Workflow https://kubernetes.io/images/docs/kubectl_drain.svg

Examples:

```
# Drain node "foo", even if there are pods not managed by a replication controller,
replica set, job, daemon set or stateful set on it
kubectl drain foo --force
```

```
# As above, but abort if there are pods not managed by a replication controller, replica
set, job, daemon set or stateful set, and use a grace period of 15 minutes
kubectl drain foo --grace-period=900
```

Options:

```
--chunk-size=500:
    Return large lists in chunks rather than all at once. Pass 0 to
    disable. This flag is beta and may change in the future.

--delete-emptydir-data=false:
    Continue even if there are pods using emptyDir (local data that will
    be deleted when the node is drained).
```

```

--disable-eviction=false:
    Force drain to use delete, even if eviction is supported. This will
    bypass checking PodDisruptionBudgets, use with caution.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

--force=false:
    Continue even if there are pods that do not declare a controller.

--grace-period=-1:
    Period of time in seconds given to each pod to terminate gracefully.
    If negative, the default value specified in the pod will be used.

--ignore-daemonsets=false:
    Ignore DaemonSet-managed pods.

--pod-selector='':
    Label selector to filter pods on the node

-l, --selector='':
    Selector (label query) to filter on, supports '=', '==', and
    '!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy
    all of the specified label constraints.

--skip-wait-for-delete-timeout=0:
    If pod DeletionTimestamp older than N seconds, skip waiting for the
    pod. Seconds must be greater than 0 to skip.

--timeout=0s:
    The length of time to wait before giving up, zero means infinite

```

Usage:

```
kubect1 drain NODE [options]
```

=====

```
kubect1 taint --help
```

Update the taints on one or more nodes.

- * A taint consists of a key, value, and effect. As an argument here, it is expressed as key=value:effect.
- * The key must begin with a letter or number, and may contain letters, numbers, hyphens, dots, and underscores, up to 253 characters.
- * Optionally, the key can begin with a DNS subdomain prefix and a single '/', like example.com/my-app.
- * The value is optional. If given, it must begin with a letter or number, and may contain letters, numbers, hyphens, dots, and underscores, up to 63 characters.
- * The effect must be NoSchedule, PreferNoSchedule or NoExecute.
- * Currently taint can only apply to node.

Examples:

```

# Update node 'foo' with a taint with key 'dedicated' and value 'special-user' and effect 'NoSchedule'
# If a taint with that key and effect already exists, its value is replaced as specified
kubect1 taint nodes foo dedicated=special-user:NoSchedule

# Remove from node 'foo' the taint with key 'dedicated' and effect 'NoSchedule' if one exists
kubect1 taint nodes foo dedicated:NoSchedule-

```

```
# Remove from node 'foo' all the taints with key 'dedicated'
kubectl taint nodes foo dedicated-

# Add a taint with key 'dedicated' on nodes having label mylabel=X
kubectl taint node -l myLabel=X dedicated=foo:PreferNoSchedule

# Add to node 'foo' a taint with key 'bar' and no value
kubectl taint nodes foo bar:NoSchedule
```

Options:

```
--all=false:
    Select all nodes in the cluster

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go lang and jsonpath output
    formats.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

--field-manager='kubectl-taint':
    Name of the manager used to track field ownership.

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file).

--overwrite=false:
    If true, allow taints to be overwritten, otherwise reject taint
    updates that overwrite existing taints.

-l, --selector='':
    Selector (label query) to filter on, supports '=', '==', and
    '!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy
    all of the specified label constraints.

--show-managed-fields=false:
    If true, keep the managedFields when printing objects in JSON or YAML
    format.

--template='':
    Template string or path to template file to use when -o=go-template,
    -o=go-template-file. The template format is go lang templates
    [http://golang.org/pkg/text/template/#pkg-overview].

--validate='strict':
    Must be one of: strict (or true), warn, ignore (or false).
    "true" or
    "strict" will use a schema to validate the input and fail the request
    if invalid. It will perform server side validation if
    ServerSideFieldValidation is enabled on the api-server, but will fall
    back to less reliable client-side validation if not.
    "warn" will
    warn about unknown or duplicate fields without blocking the request if
    server-side field validation is enabled on the API server, and behave
    as "ignore" otherwise.
    "false" or "ignore" will not perform any
    schema validation, silently dropping any unknown or duplicate fields.
```

Usage:

```
kubectl taint NODE NAME KEY_1=VAL_1:TAINT_EFFECT_1 ... KEY_N=VAL_N:TAINT_EFFECT_N
[options]
```



```
=====
kubectl describe --help
Show details of a specific resource or group of resources.
```

Print a detailed description of the selected resources, including related resources such as events or controllers. You may select a single object by name, all objects of that type, provide a name prefix, or label selector. For example:

```
$ kubectl describe TYPE NAME_PREFIX
```

will first check for an exact match on TYPE and NAME_PREFIX. If no such resource exists, it will output details for every resource that has a name prefixed with NAME_PREFIX.

Use "kubectl api-resources" for a complete list of supported resources.

Examples:

```
# Describe a node
kubectl describe nodes kubernetes-node-emt8.c.myproject.internal

# Describe a pod
kubectl describe pods/nginx

# Describe a pod identified by type and name in "pod.json"
kubectl describe -f pod.json

# Describe all pods
kubectl describe pods

# Describe pods by label name=myLabel
kubectl describe po -l name=myLabel

# Describe all pods managed by the 'frontend' replication controller
# (rc-created pods get the name of the rc as a prefix in the pod name)
kubectl describe pods frontend
```

Options:

```
-A, --all-namespaces=false:
    If present, list the requested object(s) across all namespaces.
    Namespace in current context is ignored even if specified with
    --namespace.

--chunk-size=500:
    Return large lists in chunks rather than all at once. Pass 0 to
    disable. This flag is beta and may change in the future.

-f, --filename=[ ]:
    Filename, directory, or URL to files containing the resource to
    describe

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
    with -f or -R.

-R, --recursive=false:
    Process the directory used in -f, --filename recursively. Useful when
    you want to manage related manifests organized within the same
    directory.

-l, --selector='':
    Selector (label query) to filter on, supports '=', '==', and
    '!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy
    all of the specified label constraints.
```

--show-events=true:
If true, display events related to the described object.

Usage:

kubectl describe (-f FILENAME | TYPE [NAME_PREFIX | -l label] | TYPE/NAME)
[options]

=====

kubectl logs --help

Print the logs for a container in a pod or specified resource. If the pod has only one container, the container name is optional.

Examples:

Return snapshot logs from pod nginx with only one container

kubectl logs nginx

Return snapshot logs from pod nginx with multi containers

kubectl logs nginx --all-containers=true

Return snapshot logs from all containers in pods defined by label app=nginx

kubectl logs -l app=nginx --all-containers=true

Return snapshot of previous terminated ruby container logs from pod web-1

kubectl logs -p -c ruby web-1

Begin streaming the logs of the ruby container in pod web-1

kubectl logs -f -c ruby web-1

Begin streaming the logs from all containers in pods defined by label app=nginx

kubectl logs -f -l app=nginx --all-containers=true

Display only the most recent 20 lines of output in pod nginx

kubectl logs --tail=20 nginx

Show all logs from pod nginx written in the last hour

kubectl logs --since=1h nginx

Show logs from a kubelet with an expired serving certificate

kubectl logs --insecure-skip-tls-verify-backend nginx

Return snapshot logs from first container of a job named hello

kubectl logs job/hello

Return snapshot logs from container nginx-1 of a deployment named nginx

kubectl logs deployment/nginx -c nginx-1

Options:

--all-containers=false:

Get all containers' logs in the pod(s).

-c, --container='':

Print the logs of this container

-f, --follow=false:

Specify if the logs should be streamed.

--ignore-errors=false:

If watching / following pod logs, allow for any errors that occur to be non-fatal

--insecure-skip-tls-verify-backend=false:

Skip verifying the identity of the kubelet that logs are requested from. In theory, an attacker could provide invalid log content back.

You might want to use this if your kubelet serving certificates have expired.

`--limit-bytes=0:`
Maximum bytes of logs to return. Defaults to no limit.

`--max-log-requests=5:`
Specify maximum number of concurrent logs to follow when using by a selector. Defaults to 5.

`--pod-running-timeout=20s:`
The length of time (like 5s, 2m, or 3h, higher than zero) to wait until at least one pod is running

`--prefix=false:`
Prefix each log line with the log source (pod name and container name)

`-p, --previous=false:`
If true, print the logs for the previous instance of the container in a pod if it exists.

`-l, --selector='':`
Selector (label query) to filter on, supports '=', '==', and '!='. (e.g. `-l key1=value1,key2=value2`). Matching objects must satisfy all of the specified label constraints.

`--since=0s:`
Only return logs newer than a relative duration like 5s, 2m, or 3h. Defaults to all logs. Only one of `since-time` / `since` may be used.

`--since-time='':`
Only return logs after a specific date (RFC3339). Defaults to all logs. Only one of `since-time` / `since` may be used.

`--tail=-1:`
Lines of recent log file to display. Defaults to -1 with no selector, showing all log lines otherwise 10, if a selector is provided.

`--timestamps=false:`
Include timestamps on each line in the log output

Usage:

`kubecttl logs [-f] [-p] (POD | TYPE/NAME) [-c CONTAINER] [options]`

=====

`kubecttl attach --help`

Attach to a process that is already running inside an existing container.

Examples:

Get output from running pod mypod; use the 'kubecttl.kubernetes.io/default-container' annotation

for selecting the container to be attached or the first container in the pod will be chosen

`kubecttl attach mypod`

Get output from ruby-container from pod mypod

`kubecttl attach mypod -c ruby-container`

Switch to raw terminal mode; sends stdin to 'bash' in ruby-container from pod mypod

and sends stdout/stderr from 'bash' back to the client

`kubecttl attach mypod -c ruby-container -i -t`

Get output from the first pod of a replica set named nginx

`kubecttl attach rs/nginx`

Options:

-c, --container='':
Container name. If omitted, use the
kubectrl.kubernetes.io/default-container annotation for selecting the
container to be attached or the first container in the pod will be
chosen

--pod-running-timeout=1m0s:
The length of time (like 5s, 2m, or 3h, higher than zero) to wait
until at least one pod is running

-q, --quiet=false:
Only print output from the remote session

-i, --stdin=false:
Pass stdin to the container

-t, --tty=false:
Stdin is a TTY

Usage:

kubectrl attach (POD | TYPE/NAME) -c CONTAINER [options]

=====

kubectrl exec --help
Execute a command in a container.

Examples:

Get output from running the 'date' command from pod mypod, using the first-container by default
kubectrl exec mypod -- date

Get output from running the 'date' command in ruby-container from pod mypod
kubectrl exec mypod -c ruby-container -- date

Switch to raw terminal mode; sends stdin to 'bash' in ruby-container from pod mypod
and sends stdout/stderr from 'bash' back to the client
kubectrl exec mypod -c ruby-container -i -t -- bash -il

List contents of /usr from the first container of pod mypod and sort by modification time
If the command you want to execute in the pod has any flags in common (e.g. -i),
you must use two dashes (--) to separate your command's flags/arguments
Also note, do not surround your command and its flags/arguments with quotes
unless that is how you would execute it normally (i.e., do ls -t /usr, not "ls -t /usr")
kubectrl exec mypod -i -t -- ls -t /usr

Get output from running 'date' command from the first pod of the deployment mydeployment, using the first container by default
kubectrl exec deploy/mydeployment -- date

Get output from running 'date' command from the first pod of the service myservice, using the first container by default
kubectrl exec svc/myservice -- date

Options:

-c, --container='':
Container name. If omitted, use the
kubectrl.kubernetes.io/default-container annotation for selecting the
container to be attached or the first container in the pod will be
chosen

-f, --filename=[]:

to use to exec into the resource

`--pod-running-timeout=1m0s:`

The length of time (like 5s, 2m, or 3h, higher than zero) to wait until at least one pod is running

`-q, --quiet=false:`

Only print output from the remote session

`-i, --stdin=false:`

Pass stdin to the container

`-t, --tty=false:`

Stdin is a TTY

Usage:

`kubectrl exec (POD | TYPE/NAME) [-c CONTAINER] [flags] -- COMMAND [args...]`
[options]

=====

`kubectrl port-forward --help`

Forward one or more local ports to a pod.

Use resource type/name such as deployment/mydeployment to select a pod.
Resource type defaults to 'pod' if omitted.

If there are multiple pods matching the criteria, a pod will be selected automatically. The forwarding session ends when the selected pod terminates, and a rerun of the command is needed to resume forwarding.

Examples:

Listen on ports 5000 and 6000 locally, forwarding data to/from ports 5000 and 6000 in the pod

`kubectrl port-forward pod/mypod 5000 6000`

Listen on ports 5000 and 6000 locally, forwarding data to/from ports 5000 and 6000 in a pod selected by the deployment

`kubectrl port-forward deployment/mydeployment 5000 6000`

Listen on port 8443 locally, forwarding to the targetPort of the service's port named "https" in a pod selected by the service

`kubectrl port-forward service/myservice 8443:https`

Listen on port 8888 locally, forwarding to 5000 in the pod

`kubectrl port-forward pod/mypod 8888:5000`

Listen on port 8888 on all addresses, forwarding to 5000 in the pod

`kubectrl port-forward --address 0.0.0.0 pod/mypod 8888:5000`

Listen on port 8888 on localhost and selected IP, forwarding to 5000 in the pod

`kubectrl port-forward --address localhost,10.19.21.23 pod/mypod 8888:5000`

Listen on a random port locally, forwarding to 5000 in the pod

`kubectrl port-forward pod/mypod :5000`

Options:

`--address=[localhost]:`

Addresses to listen on (comma separated). Only accepts IP addresses or localhost as a value. When localhost is supplied, kubectrl will try to bind on both 127.0.0.1 and ::1 and will fail if neither of these addresses are available to bind.

`--pod-running-timeout=1m0s:`

The length of time (like 5s, 2m, or 3h, higher than zero) to wait until at least one pod is running

Usage:

```
kubectl port-forward TYPE/NAME [options] [LOCAL_PORT:]REMOTE_PORT [...  
[LOCAL_PORT_N:]REMOTE_PORT_N]
```

```
=====
```

```
kubectl proxy --help
```

Creates a proxy server or application-level gateway between localhost and the Kubernetes API server. It also allows serving static content over specified HTTP path. All incoming data enters through one port and gets forwarded to the remote Kubernetes API server port, except for the path matching the static content path.

Examples:

```
# To proxy all of the Kubernetes API and nothing else  
kubectl proxy --api-prefix=/  

```

```
# To proxy only part of the Kubernetes API and also some static files  
# You can get pods info with 'curl localhost:8001/api/v1/pods'  
kubectl proxy --www=/my/files --www-prefix=/static/ --api-prefix=/api/  

```

```
# To proxy the entire Kubernetes API at a different root  
# You can get pods info with 'curl localhost:8001/custom/api/v1/pods'  
kubectl proxy --api-prefix=/custom/  

```

```
# Run a proxy to the Kubernetes API server on port 8011, serving static content from  
./local/www/  
kubectl proxy --port=8011 --www=./local/www/  

```

```
# Run a proxy to the Kubernetes API server on an arbitrary local port  
# The chosen port for the server will be output to stdout  
kubectl proxy --port=0  

```

```
# Run a proxy to the Kubernetes API server, changing the API prefix to k8s-api  
# This makes e.g. the pods API available at localhost:8001/k8s-api/v1/pods/  
kubectl proxy --api-prefix=/k8s-api  

```

Options:

```
--accept-hosts='^localhost$,^127\.0\.0\.1$,^\[:1\]$':  
    Regular expression for hosts that the proxy should accept.
```

```
--accept-paths='^.*':  
    Regular expression for paths that the proxy should accept.
```

```
--address='127.0.0.1':  
    The IP address on which to serve on.
```

```
--api-prefix='/':  
    Prefix to serve the proxied API under.
```

```
--append-server-path=false:  
    If true, enables automatic path appending of the kube context server  
    path to each request.
```

```
--disable-filter=false:  
    If true, disable request filtering in the proxy. This is dangerous,  
    and can leave you vulnerable to XSRF attacks, when used with an  
    accessible port.
```

```
--keepalive=0s:  
    keepalive specifies the keep-alive period for an active network  
    connection. Set to 0 to disable keepalive.
```

```
-p, --port=8001:  
    The port on which to run the proxy. Set to 0 to pick a random port.
```

```

--reject-methods='^$':
    Regular expression for HTTP methods that the proxy should reject
    (example --reject-methods='POST,PUT,PATCH').

--reject-paths='^/api/.*/*pods/.*/*exec,^/api/.*/*pods/.*/*attach':
    Regular expression for paths that the proxy should reject. Paths
    specified here will be rejected even accepted by --accept-paths.

-u, --unix-socket='':
    Unix socket on which to run the proxy.

-w, --www='':
    Also serve static files from the given directory under the specified
    prefix.

-P, --www-prefix='/static/':
    Prefix to serve static files under, if static file directory is
    specified.

```

Usage:

```

kubectrl proxy [--port=PORT] [--www=static-dir] [--www-prefix=prefix] [--api-prefix=prefix]
[options]

```

```

=====
kubectrl cp --help
Copy files and directories to and from containers.

```

Examples:

```

# !!!Important Note!!!
# Requires that the 'tar' binary is present in your container
# image. If 'tar' is not present, 'kubectrl cp' will fail.
#
# For advanced use cases, such as symlinks, wildcard expansion or
# file mode preservation, consider using 'kubectrl exec'.

# Copy /tmp/foo local file to /tmp/bar in a remote pod in namespace
<some-namespace>
tar cf - /tmp/foo | kubectrl exec -i -n <some-namespace> <some-pod> -- tar xf - -C /tmp/bar

# Copy /tmp/foo from a remote pod to /tmp/bar locally
kubectrl exec -n <some-namespace> <some-pod> -- tar cf - /tmp/foo | tar xf - -C /tmp/bar

# Copy /tmp/foo_dir local directory to /tmp/bar_dir in a remote pod in the default
namespace
kubectrl cp /tmp/foo_dir <some-pod>:/tmp/bar_dir

# Copy /tmp/foo local file to /tmp/bar in a remote pod in a specific container
kubectrl cp /tmp/foo <some-pod>:/tmp/bar -c <specific-container>

# Copy /tmp/foo local file to /tmp/bar in a remote pod in namespace <some-namespace>
kubectrl cp /tmp/foo <some-namespace>/<some-pod>:/tmp/bar

# Copy /tmp/foo from a remote pod to /tmp/bar locally
kubectrl cp <some-namespace>/<some-pod>:/tmp/foo /tmp/bar

```

Options:

```

-c, --container='':
    Container name. If omitted, use the
    kubectrl.kubernetes.io/default-container annotation for selecting the
    container to be attached or the first container in the pod will be
    chosen

--no-preserve=false:
    The copied file/directory's ownership and permissions will not be

```

preserved in the container

--retries=0:

Set number of retries to complete a copy operation from a container.
Specify 0 to disable or any negative value for infinite retrying. The
default is 0 (no retry).

Usage:

kubectl cp <file-spec-src> <file-spec-dest> [options]

=====

kubectl auth --help
Inspect authorization

Available Commands:

can-i Check whether an action is allowed
reconcile Reconciles rules for RBAC role, role binding, cluster role, and
cluster role binding objects

Usage:

kubectl auth [flags] [options]

=====

kubectl debug --help
Debug cluster resources using interactive debugging containers.

'debug' provides automation for common debugging tasks for cluster objects
identified by resource and name. Pods will be used by default if no resource is
specified.

The action taken by 'debug' varies depending on what resource is specified.
Supported actions include:

- * Workload: Create a copy of an existing pod with certain attributes changed,
for example changing the image tag to a new version.
- * Workload: Add an ephemeral container to an already running pod, for example
to add debugging utilities without restarting the pod.
- * Node: Create a new pod that runs in the node's host namespaces and can
access the node's filesystem.

Examples:

```
# Create an interactive debugging session in pod mypod and immediately attach to it.
# (requires the EphemeralContainers feature to be enabled in the cluster)
kubectl debug mypod -it --image=busybox

# Create a debug container named debugger using a custom automated debugging image.
# (requires the EphemeralContainers feature to be enabled in the cluster)
kubectl debug --image=myproj/debug-tools -c debugger mypod

# Create a copy of mypod adding a debug container and attach to it
kubectl debug mypod -it --image=busybox --copy-to=my-debugger

# Create a copy of mypod changing the command of mycontainer
kubectl debug mypod -it --copy-to=my-debugger --container=mycontainer -- sh

# Create a copy of mypod changing all container images to busybox
kubectl debug mypod --copy-to=my-debugger --set-image *=busybox

# Create a copy of mypod adding a debug container and changing container images
kubectl debug mypod -it --copy-to=my-debugger --image=debian --set-
image=app=app:debug,sidecar=sidecar:debug
```



```
# Create an interactive debugging session on a node and immediately attach to it.
# The container will run in the host namespaces and the host's filesystem will be mounted
at /host
kubectl debug node/mynode -it --image=busybox
```

Options:

```
--arguments-only=false:
    If specified, everything after -- will be passed to the new container
    as Args instead of Command.

--attach=false:
    If true, wait for the container to start running, and then attach as
    if 'kubectl attach ...' were called. Default false, unless
    '-i/--stdin' is set, in which case the default is true.

-c, --container='':
    Container name to use for debug container.

--copy-to='':
    Create a copy of the target Pod with this name.

--env=[ ]:
    Environment variables to set in the container.

--image='':
    Container image to use for debug container.

--image-pull-policy='':
    The image pull policy for the container. If left empty, this value
    will not be specified by the client and defaulted by the server.

--profile='legacy':
    Debugging profile.

-q, --quiet=false:
    If true, suppress informational messages.

--replace=false:
    When used with '--copy-to', delete the original Pod.

--same-node=false:
    When used with '--copy-to', schedule the copy of target Pod on the
    same node.

--set-image=[ ]:
    When used with '--copy-to', a list of name=image pairs for changing
    container images, similar to how 'kubectl set image' works.

--share-processes=true:
    When used with '--copy-to', enable process namespace sharing in the
    copy.

-i, --stdin=false:
    Keep stdin open on the container(s) in the pod, even if nothing is
    attached.

--target='':
    When using an ephemeral container, target processes in this container
    name.

-t, --tty=false:
    Allocate a TTY for the debugging container.
```

Usage:

```
kubectl debug (POD | TYPE[.VERSION].GROUP)/NAME [ -- COMMAND [args...] ] [options]
```

```
=====
kubecttl events --help
Display events
```

Prints a table of the most important information about events. You can request events for a namespace, for all namespace, or filtered to only those pertaining to a specified resource.

Examples:

```
# List recent events in the default namespace.
kubecttl events

# List recent events in all namespaces.
kubecttl events --all-namespaces

# List recent events for the specified pod, then wait for more events and list them as
they arrive.
kubecttl events --for pod/web-pod-13je7 --watch

# List recent events in given format. Supported ones, apart from default, are json and
yaml.
kubecttl events -oyaml

# List recent only events in given event types
kubecttl events --types=Warning,Normal
```

Options:

```
-A, --all-namespaces=false:
    If present, list the requested object(s) across all namespaces.
    Namespace in current context is ignored even if specified with
    --namespace.

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go lang and jsonpath output
    formats.

--chunk-size=500:
    Return large lists in chunks rather than all at once. Pass 0 to
    disable. This flag is beta and may change in the future.

--for='':
    Filter events to only those pertaining to the specified resource.

--no-headers=false:
    When using the default output format, don't print headers.

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file).

--show-managed-fields=false:
    If true, keep the managedFields when printing objects in JSON or YAML
    format.

--template='':
    Template string or path to template file to use when -o=go-template,
    -o=go-template-file. The template format is go lang templates
    [http://golang.org/pkg/text/template/#pkg-overview].

--types=[ ]:
    Output only events of given types.

-w, --watch=false:
```

After listing the requested events, watch for more events.

Usage:

```
kubectl events
[(-o|--output=)json|yaml|name|go-template|go-template-file|template|templatefile|jsonpath|
jsonpath-as-json|jsonpath-file]
[--for TYPE/NAME] [--watch] [--event=Normal,Warning] [options]
```

=====

`kubectl diff --help`

Diff configurations specified by file name or stdin between the current online configuration, and the configuration as it would be if applied.

The output is always YAML.

KUBECTL_EXTERNAL_DIFF environment variable can be used to select your own diff command. Users can use external commands with params too, example:
KUBECTL_EXTERNAL_DIFF="colordiff -N -u"

By default, the "diff" command available in your path will be run with the "-u" (unified diff) and "-N" (treat absent files as empty) options.

Exit status: 0 No differences were found. 1 Differences were found. >1 Kubectl or diff failed with an error.

Note: KUBECTL_EXTERNAL_DIFF, if used, is expected to follow that convention.

Examples:

```
# Diff resources included in pod.json
kubectl diff -f pod.json

# Diff file read from stdin
cat service.yaml | kubectl diff -f -
```

Options:

```
--field-manager='kubectl-client-side-apply':
    Name of the manager used to track field ownership.

-f, --filename=[ ]:
    Filename, directory, or URL to files contains the configuration to
    diff

--force-conflicts=false:
    If true, server-side apply will force the changes against conflicts.

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
    with -f or -R.

--prune=false:
    Include resources that would be deleted by pruning. Can be used with
    -l and default shows all resources would be pruned

--prune-allowlist=[ ]:
    Overwrite the default whitelist with <group/version/kind> for --prune

-R, --recursive=false:
    Process the directory used in -f, --filename recursively. Useful when
    you want to manage related manifests organized within the same
    directory.

-l, --selector='':
    Selector (label query) to filter on, supports '=', '==', and
```

'!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy all of the specified label constraints.

--server-side=false:

If true, apply runs in the server instead of the client.

--show-managed-fields=false:

If true, include managed fields in the diff.

Usage:

kubectl diff -f FILENAME [options]

=====

kubectl apply --help

Apply a configuration to a resource by file name or stdin. The resource name must be specified. This resource will be created if it doesn't exist yet. To use 'apply', always create the resource initially with either 'apply' or 'create --save-config'.

JSON and YAML formats are accepted.

Alpha Disclaimer: the --prune functionality is not yet complete. Do not use unless you are aware of what the current state is. See <https://issues.k8s.io/34274>.

Examples:

Apply the configuration in pod.json to a pod
kubectl apply -f ./pod.json

Apply resources from a directory containing kustomization.yaml - e.g.
dir/kustomization.yaml
kubectl apply -k dir/

Apply the JSON passed into stdin to a pod
cat pod.json | kubectl apply -f -

Apply the configuration from all files that end with '.json' - i.e. expand wildcard characters in file names
kubectl apply -f '*.json'

Note: --prune is still in Alpha
Apply the configuration in manifest.yaml that matches label app=nginx and delete all other resources that are not in the file and match label app=nginx
kubectl apply --prune -f manifest.yaml -l app=nginx

Apply the configuration in manifest.yaml and delete all the other config maps that are not in the file
kubectl apply --prune -f manifest.yaml --all --prune-allowlist=core/v1/ConfigMap

Available Commands:

edit-last-applied Edit latest last-applied-configuration annotations of are source/object
set-last-applied Set the last-applied-configuration annotation on a live object to match the contents of a file
view-last-applied View the latest last-applied-configuration annotations of a resource/object

Options:

--all=false:

Select all resources in the namespace of the specified resource types.

--allow-missing-template-keys=true:

If true, ignore any errors in templates when a field or map key is missing in the template. Only applies to go lang and jsonpath output

formats.

`--cascade='background':`
Must be "background", "orphan", or "foreground". Selects the deletion cascading strategy for the dependents (e.g. Pods created by a ReplicationController). Defaults to background.

`--dry-run='none':`
Must be "none", "server", or "client". If client strategy, only print the object that would be sent, without sending it. If server strategy, submit server-side request without persisting the resource.

`--field-manager='kubectl-client-side-apply':`
Name of the manager used to track field ownership.

`-f, --filename=[]:`
The files that contain the configurations to apply.

`--force=false:`
If true, immediately remove resources from API and bypass graceful deletion. Note that immediate deletion of some resources may result in inconsistency or data loss and requires confirmation.

`--force-conflicts=false:`
If true, server-side apply will force the changes against conflicts.

`--grace-period=-1:`
Period of time in seconds given to the resource to terminate gracefully. Ignored if negative. Set to 1 for immediate shutdown. Can only be set to 0 when `--force` is true (force deletion).

`-k, --kustomize='':`
Process a kustomization directory. This flag can't be used together with `-f` or `-R`.

`--openapi-patch=true:`
If true, use openapi to calculate diff when the openapi presents and the resource can be found in the openapi spec. Otherwise, fall back to use baked-in types.

`-o, --output='':`
Output format. One of: (json, yaml, name, go-template, go-template-file, template, templatefile, jsonpath, jsonpath-as-json, jsonpath-file).

`--overwrite=true:`
Automatically resolve conflicts between the modified and live configuration by using values from the modified configuration

`--prune=false:`
Automatically delete resource objects, that do not appear in the configs and are created by either apply or create `--save-config`. Should be used with either `-l` or `--all`.

`--prune-allowlist=[]:`
Overwrite the default allowlist with <group/version/kind> for `--prune`

`-R, --recursive=false:`
Process the directory used in `-f, --filename` recursively. Useful when you want to manage related manifests organized within the same directory.

`-l, --selector='':`
Selector (label query) to filter on, supports '=', '==', and '!='. (e.g. `-l key1=value1,key2=value2`). Matching objects must satisfy all of the specified label constraints.

```
--server-side=false:
    If true, apply runs in the server instead of the client.

--show-managed-fields=false:
    If true, keep the managedFields when printing objects in JSON or YAML
    format.

--template='':
    Template string or path to template file to use when -o=go-template,
    -o=go-template-file. The template format is go lang templates
    [http://golang.org/pkg/text/template/#pkg-overview].

--timeout=0s:
    The length of time to wait before giving up on a delete, zero means
    determine a timeout from the size of the object

--validate='strict':
    Must be one of: strict (or true), warn, ignore (or false).
        "true" or
        "strict" will use a schema to validate the input and fail the request
        if invalid. It will perform server side validation if
        ServerSideFieldValidation is enabled on the api-server, but will fall
        back to less reliable client-side validation if not.
        "warn" will
        warn about unknown or duplicate fields without blocking the request if
        server-side field validation is enabled on the API server, and behave
        as "ignore" otherwise.
        "false" or "ignore" will not perform any
        schema validation, silently dropping any unknown or duplicate fields.

--wait=false:
    If true, wait for resources to be gone before returning. This waits
    for finalizers.
```

Usage:

```
kubectl apply (-f FILENAME | -k DIRECTORY) [options]
```

```
kubectl patch --help
```

Update fields of a resource using strategic merge patch, a JSON merge patch, or a JSON patch.

JSON and YAML formats are accepted.

Note: Strategic merge patch is not supported for custom resources.

Examples:

```
# Partially update a node using a strategic merge patch, specifying the patch as JSON
kubectl patch node k8s-node-1 -p '{"spec":{"unschedulable":true}}'
```

```
# Partially update a node using a strategic merge patch, specifying the patch as YAML
kubectl patch node k8s-node-1 -p $'spec:\n unschedulable: true'
```

```
# Partially update a node identified by the type and name specified in "node.json" using
strategic merge patch
```

```
kubectl patch -f node.json -p '{"spec":{"unschedulable":true}}'
```

```
# Update a container's image; spec.containers[*].name is required because it's a merge key
kubectl patch pod valid-pod -p
```

```
'{"spec":{"containers":[{"name":"kubernetes-serve-hostname","image":"newimage"}]}}'
```

```
# Update a container's image using a JSON patch with positional arrays
kubectl patch pod valid-pod --type='json' -p='[{"op": "replace", "path":
"/spec/containers/0/image", "value":"new image"}]'
```

```
# Update a deployment's replicas through the scale subresource using a merge patch.
kubectl patch deployment nginx-deployment --subresource='scale' --type='merge' -p
'{"spec":{"replicas":2}}'
```

Options:

```
--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to goyaml and jsonpath output
    formats.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

--field-manager='kubectl-patch':
    Name of the manager used to track field ownership.

-f, --filename=[ ]:
    Filename, directory, or URL to files identifying the resource to
    update

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
    with -f or -R.

--local=false:
    If true, patch will operate on the content of the file, not the
    server-side resource.

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file).

-p, --patch='':
    The patch to be applied to the resource JSON file.

--patch-file='':
    A file containing a patch to be applied to the resource.

-R, --recursive=false:
    Process the directory used in -f, --filename recursively. Useful when
    you want to manage related manifests organized within the same
    directory.

--show-managed-fields=false:
    If true, keep the managedFields when printing objects in JSON or YAML
    format.

--subresource='':
    If specified, patch will operate on the subresource of the requested
    object. Must be one of [status scale]. This flag is alpha and may
    change in the future.

--template='':
    Template string or path to template file to use when -o=go-template,
    -o=go-template-file. The template format is goyaml templates
    [http://golang.org/pkg/text/template/#pkg-overview].

--type='strategic':
    The type of patch being provided; one of [json merge strategic]
```

Usage:

```
kubectl patch (-f FILENAME | TYPE NAME) [-p PATCH|--patch-file FILE] [options]
```

```
=====
kubectl replace --help
Replace a resource by file name or stdin.
```

JSON and YAML formats are accepted. If replacing an existing resource, the complete resource spec must be provided. This can be obtained by

```
$ kubectl get TYPE NAME -o yaml
```

Examples:

```
# Replace a pod using the data in pod.json
kubectl replace -f ./pod.json
```

```
# Replace a pod based on the JSON passed into stdin
cat pod.json | kubectl replace -f -
```

```
# Update a single-container pod's image version (tag) to v4
kubectl get pod mypod -o yaml | sed 's/\\(image: myimage\\):.*$/\\1:v4/' | kubectl replace -f -
```

```
# Force replace, delete and then re-create the resource
kubectl replace --force -f ./pod.json
```

Options:

--allow-missing-template-keys=true:

If true, ignore any errors in templates when a field or map key is missing in the template. Only applies to go lang and jsonpath output formats.

--cascade='background':

Must be "background", "orphan", or "foreground". Selects the deletion cascading strategy for the dependents (e.g. Pods created by a ReplicationController). Defaults to background.

--dry-run='none':

Must be "none", "server", or "client". If client strategy, only print the object that would be sent, without sending it. If server strategy, submit server-side request without persisting the resource.

--field-manager='kubectl-replace':

Name of the manager used to track field ownership.

-f, --filename=[]:

The files that contain the configurations to replace.

--force=false:

If true, immediately remove resources from API and bypass graceful deletion. Note that immediate deletion of some resources may result in inconsistency or data loss and requires confirmation.

--grace-period=-1:

Period of time in seconds given to the resource to terminate gracefully. Ignored if negative. Set to 1 for immediate shutdown. Can only be set to 0 when --force is true (force deletion).

-k, --kustomize='':

Process a kustomization directory. This flag can't be used together with -f or -R.

-o, --output='':

Output format. One of: (json, yaml, name, go-template, go-template-file, template, templatefile, jsonpath, jsonpath-as-json, jsonpath-file).

--raw='':
Raw URI to PUT to the server. Uses the transport specified by the kubeconfig file.

-R, --recursive=false:
Process the directory used in -f, --filename recursively. Useful when you want to manage related manifests organized within the same directory.

--save-config=false:
If true, the configuration of current object will be saved in its annotation. Otherwise, the annotation will be unchanged. This flag is useful when you want to perform kubectl apply on this object in the future.

--show-managed-fields=false:
If true, keep the managedFields when printing objects in JSON or YAML format.

--subresource='':
If specified, replace will operate on the subresource of the requested object. Must be one of [status scale]. This flag is alpha and may change in the future.

--template='':
Template string or path to template file to use when -o=go-template, -o=go-template-file. The template format is golang templates [<http://golang.org/pkg/text/template/#pkg-overview>].

--timeout=0s:
The length of time to wait before giving up on a delete, zero means determine a timeout from the size of the object

--validate='strict':
Must be one of: strict (or true), warn, ignore (or false).
"true" or "strict" will use a schema to validate the input and fail the request if invalid. It will perform server side validation if ServerSideFieldValidation is enabled on the api-server, but will fall back to less reliable client-side validation if not.
"warn" will warn about unknown or duplicate fields without blocking the request if server-side field validation is enabled on the API server, and behave as "ignore" otherwise.
"false" or "ignore" will not perform any schema validation, silently dropping any unknown or duplicate fields.

--wait=false:
If true, wait for resources to be gone before returning. This waits for finalizers.

Usage:

```
kubectl replace -f FILENAME [options]
```

```
kubectl wait --help
```

Experimental: Wait for a specific condition on one or many resources.

The command takes multiple resources and waits until the specified condition is seen in the Status field of every given resource.

Alternatively, the command can wait for the given set of resources to be deleted by providing the "delete" keyword as the value to the --for flag.

A successful message will be printed to stdout indicating when the specified condition has been met. You can use -o option to change to output destination.

Examples:

```
# Wait for the pod "busybox1" to contain the status condition of type "Ready"
kubectl wait --for=condition=Ready pod/busybox1

# The default value of status condition is true; you can wait for other targets after an
equal delimiter (compared after Unicode simple case folding, which is a more general form of
case-insensitivity):
kubectl wait --for=condition=Ready=false pod/busybox1

# Wait for the pod "busybox1" to contain the status phase to be "Running".
kubectl wait --for=jsonpath='{.status.phase}'=Running pod/busybox1

# Wait for the pod "busybox1" to be deleted, with a timeout of 60s, after having issued
the "delete" command
kubectl delete pod/busybox1
kubectl wait --for=delete pod/busybox1 --timeout=60s
```

Options:

```
--all=false:
    Select all resources in the namespace of the specified resource types

-A, --all-namespaces=false:
    If present, list the requested object(s) across all namespaces.
    Namespace in current context is ignored even if specified with
    --namespace.

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go lang and jsonpath output
    formats.

--field-selector='':
    Selector (field query) to filter on, supports '=', '==', and
    '!='. (e.g. --field-selector key1=value1,key2=value2). The server only
    supports a limited number of field queries per type.

-f, --filename=[ ]:
    identifying the resource.

--for='':
    The condition to wait on:
    [delete|condition=condition-name[=condition-value]|jsonpath='{JSONPath
    expression}'=JSONPath Condition]. The default condition-value is true.
    Condition values are compared after Unicode simple case folding, which
    is a more general form of case-insensitivity.

--local=false:
    If true, annotation will NOT contact api-server but run locally.

-o, --output='':
    Output format. One of: (json, yaml, name, go-template,
    go-template-file, template, templatefile, jsonpath, jsonpath-as-json,
    jsonpath-file).

-R, --recursive=true:
    Process the directory used in -f, --filename recursively. Useful when
    you want to manage related manifests organized within the same
    directory.

-l, --selector='':
    Selector (label query) to filter on, supports '=', '==', and
    '!='. (e.g. -l key1=value1,key2=value2)
```

--show-managed-fields=false:
If true, keep the managedFields when printing objects in JSON or YAML format.

--template='':
Template string or path to template file to use when -o=go-template, -o=go-template-file. The template format is golang templates [http://golang.org/pkg/text/template/#pkg-overview].

--timeout=30s:
The length of time to wait before giving up. Zero means check once and don't wait, negative means wait for a week.

Usage:

kubectl wait [--f FILENAME] | resource.group/resource.name | resource.group [(-l label | --all)] [--for=delete|--for condition=available|--for=jsonpath='{}'=value] [options]

=====

kubectl kustomize --help
Build a set of KRM resources using a 'kustomization.yaml' file. The DIR argument must be a path to a directory containing 'kustomization.yaml', or a git repository URL with a path suffix specifying same with respect to the repository root. If DIR is omitted, '.' is assumed.

Examples:

Build the current working directory
kubectl kustomize

Build some shared configuration directory
kubectl kustomize /home/config/production

Build from github
kubectl kustomize
<https://github.com/kubernetes-sigs/kustomize.git/examples/helloWorld?ref=v1.0.6>

Options:

--as-current-user=false:
use the uid and gid of the command executor to run the function in the container

--enable-alpha-plugins=false:
enable kustomize plugins

--enable-helm=false:
Enable use of the Helm chart inflator generator.

-e, --env=[]:
a list of environment variables to be used by functions

--helm-command='helm':
helm command (path to executable)

--load-restrictor='LoadRestrictionsRootOnly':
if set to 'LoadRestrictionsNone', local kustomizations may load files from outside their root. This does, however, break the relocatability of the kustomization.

--mount=[]:
a list of storage options read from the filesystem

--network=false:
enable network access for functions that declare it

```
--network-name='bridge':
    the docker network to run the container in

-o, --output='':
    If specified, write output to this path.

--reorder='legacy':
    Reorder the resources just before output. Use 'legacy' to apply a
    legacy reordering (Namespaces first, Webhooks last, etc). Use 'none'
    to suppress a final reordering.
```

Usage:

```
kubectl kustomize DIR [flags] [options]
```

=====

kubectl label --help

Update the labels on a resource.

- * A label key and value must begin with a letter or number, and may contain letters, numbers, hyphens, dots, and underscores, up to 63 characters each.
- * Optionally, the key can begin with a DNS subdomain prefix and a single '/', like example.com/my-app.
- * If --overwrite is true, then existing labels can be overwritten, otherwise attempting to overwrite a label will result in an error.
- * If --resource-version is specified, then updates will use this resource version, otherwise the existing resource-version will be used.

Examples:

```
# Update pod 'foo' with the label 'unhealthy' and the value 'true'
kubectl label pods foo unhealthy=true

# Update pod 'foo' with the label 'status' and the value 'unhealthy', overwriting any
existing value
kubectl label --overwrite pods foo status=unhealthy

# Update all pods in the namespace
kubectl label pods --all status=unhealthy

# Update a pod identified by the type and name in "pod.json"
kubectl label -f pod.json status=unhealthy

# Update pod 'foo' only if the resource is unchanged from version 1
kubectl label pods foo status=unhealthy --resource-version=1

# Update pod 'foo' by removing a label named 'bar' if it exists
# Does not require the --overwrite flag
kubectl label pods foo bar-
```

Options:

```
--all=false:
    Select all resources, in the namespace of the specified resource types

-A, --all-namespaces=false:
    If true, check the specified action in all namespaces.

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to go lang and jsonpath output
    formats.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.
```

`--field-manager='kubectl-label':`
 Name of the manager used to track field ownership.

`--field-selector='':`
 Selector (field query) to filter on, supports '=', '==', and '!='. (e.g. `--field-selector key1=value1,key2=value2`). The server only supports a limited number of field queries per type.

`-f, --filename=[]:`
 Filename, directory, or URL to files identifying the resource to update the labels

`-k, --kustomize='':`
 Process the kustomization directory. This flag can't be used together with `-f` or `-R`.

`--list=false:`
 If true, display the labels for a given resource.

`--local=false:`
 If true, label will NOT contact api-server but run locally.

`-o, --output='':`
 Output format. One of: (json, yaml, name, go-template, go-template-file, template, templatefile, jsonpath, jsonpath-as-json, jsonpath-file).

`--overwrite=false:`
 If true, allow labels to be overwritten, otherwise reject label updates that overwrite existing labels.

`-R, --recursive=false:`
 Process the directory used in `-f, --filename` recursively. Useful when you want to manage related manifests organized within the same directory.

`--resource-version='':`
 If non-empty, the labels update will only succeed if this is the current resource-version for the object. Only valid when specifying a single resource.

`-l, --selector='':`
 Selector (label query) to filter on, supports '=', '==', and '!='. (e.g. `-l key1=value1,key2=value2`). Matching objects must satisfy all of the specified label constraints.

`--show-managed-fields=false:`
 If true, keep the managedFields when printing objects in JSON or YAML format.

`--template='':`
 Template string or path to template file to use when `-o=go-template, -o=go-template-file`. The template format is go lang templates [<http://golang.org/pkg/text/template/#pkg-overview>].

Usage:

```
kubectl label [--overwrite] (-f FILENAME | TYPE NAME) KEY_1=VAL_1 ... KEY_N=VAL_N [--resource-version=version] [options]
```

=====

`kubectl annotate --help`
 Update the annotations on one or more resources.

All Kubernetes objects support the ability to store additional data with the object as annotations. Annotations are key/value pairs that can be larger than labels and include arbitrary string values such as structured JSON. Tools and system extensions may use annotations to store their own data.

Attempting to set an annotation that already exists will fail unless `--overwrite` is set. If `--resource-version` is specified and does not match the current resource version on the server the command will fail.

Use `"kubectl api-resources"` for a complete list of supported resources.

Examples:

```
# Update pod 'foo' with the annotation 'description' and the value 'my frontend'
# If the same annotation is set multiple times, only the last value will be applied
kubectl annotate pods foo description='my frontend'

# Update a pod identified by type and name in "pod.json"
kubectl annotate -f pod.json description='my frontend'

# Update pod 'foo' with the annotation 'description' and the value 'my frontend running
nginx', overwriting any existing value
kubectl annotate --overwrite pods foo description='my frontend running nginx'

# Update all pods in the namespace
kubectl annotate pods --all description='my frontend running nginx'

# Update pod 'foo' only if the resource is unchanged from version 1
kubectl annotate pods foo description='my frontend running nginx'
--resource-version=1

# Update pod 'foo' by removing an annotation named 'description' if it exists
# Does not require the --overwrite flag
kubectl annotate pods foo description-
```

Options:

```
--all=false:
    Select all resources, in the namespace of the specified resource
    types.

-A, --all-namespaces=false:
    If true, check the specified action in all namespaces.

--allow-missing-template-keys=true:
    If true, ignore any errors in templates when a field or map key is
    missing in the template. Only applies to goyaml and jsonpath output
    formats.

--dry-run='none':
    Must be "none", "server", or "client". If client strategy, only print
    the object that would be sent, without sending it. If server strategy,
    submit server-side request without persisting the resource.

--field-manager='kubectl-annotate':
    Name of the manager used to track field ownership.

--field-selector='':
    Selector (field query) to filter on, supports '=', '==', and
    '!='. (e.g. --field-selector key1=value1,key2=value2). The server only
    supports a limited number of field queries per type.

-f, --filename=[ ]:
    Filename, directory, or URL to files identifying the resource to
    update the annotation

-k, --kustomize='':
    Process the kustomization directory. This flag can't be used together
```

with -f or -R.

--list=false:

If true, display the annotations for a given resource.

--local=false:

If true, annotation will NOT contact api-server but run locally.

-o, --output='':

Output format. One of: (json, yaml, name, go-template, go-template-file, template, templatefile, jsonpath, jsonpath-as-json, jsonpath-file).

--overwrite=false:

If true, allow annotations to be overwritten, otherwise reject annotation updates that overwrite existing annotations.

-R, --recursive=false:

Process the directory used in -f, --filename recursively. Useful when you want to manage related manifests organized within the same directory.

--resource-version='':

If non-empty, the annotation update will only succeed if this is the current resource-version for the object. Only valid when specifying a single resource.

-l, --selector='':

Selector (label query) to filter on, supports '=', '==', and '!='. (e.g. -l key1=value1,key2=value2). Matching objects must satisfy all of the specified label constraints.

--show-managed-fields=false:

If true, keep the managedFields when printing objects in JSON or YAML format.

--template='':

Template string or path to template file to use when -o=go-template, -o=go-template-file. The template format is go lang templates [http://golang.org/pkg/text/template/#pkg-overview].

Usage:

kubectl annotate [--overwrite] (-f FILENAME | TYPE NAME) KEY_1=VAL_1 ... KEY_N=VAL_N [--resource-version=version] [options]

kubectl completion --help

Output shell completion code for the specified shell (bash, zsh, fish, or powershell). The shell code must be evaluated to provide interactive completion of kubectl commands. This can be done by sourcing it from the .bash_profile.

Detailed instructions on how to do this are available here:

for macOS:

<https://kubernetes.io/docs/tasks/tools/install-kubectl-macos/#enable-shell-autocompletion>

for linux:

<https://kubernetes.io/docs/tasks/tools/install-kubectl-linux/#enable-shell-autocompletion>

for windows:

<https://kubernetes.io/docs/tasks/tools/install-kubectl-windows/#enable-shell-autocompletion>

Note for zsh users: [1] zsh completions are only supported in versions of zsh
>= 5.2.

Examples:

```
# Installing bash completion on macOS using homebrew
## If running Bash 3.2 included with macOS
brew install bash-completion
## or, if running Bash 4.1+
brew install bash-completion@2
## If kubectl is installed via homebrew, this should start working immediately
## If you've installed via other means, you may need add the completion to your completion
directory
kubectl completion bash > $(brew --prefix)/etc/bash_completion.d/kubectl
```

```
# Installing bash completion on Linux
## If bash-completion is not installed on Linux, install the 'bash-completion' package
## via your distribution's package manager.
## Load the kubectl completion code for bash into the current shell
source <(kubectl completion bash)
## Write bash completion code to a file and source it from .bash_profile
kubectl completion bash > ~/.kube/completion.bash.inc printf "
# Kubectl shell completion
source '$HOME/.kube/completion.bash.inc'
" >> $HOME/.bash_profile
source $HOME/.bash_profile

# Load the kubectl completion code for zsh[1] into the current shell
source <(kubectl completion zsh)
# Set the kubectl completion code for zsh[1] to autoload on startup
kubectl completion zsh > "${fpath[1]}/_kubectl"
```

```
# Load the kubectl completion code for fish[2] into the current shell
kubectl completion fish | source
# To load completions for each session, execute once:
kubectl completion fish > ~/.config/fish/completions/kubectl.fish
```

```
# Load the kubectl completion code for powershell into the current shell
kubectl completion powershell | Out-String | Invoke-Expression
# Set kubectl completion code for powershell to run on startup
## Save completion code to a script and execute in the profile
kubectl completion powershell > $HOME\.kube\completion.ps1
Add-Content $PROFILE "$HOME\.kube\completion.ps1"
## Execute completion code in the profile
Add-Content $PROFILE "if (Get-Command kubectl -ErrorAction SilentlyContinue) {
kubectl completion powershell | Out-String | Invoke-Expression
}"
## Add completion code directly to the $PROFILE script
kubectl completion powershell >> $PROFILE
```

Usage:

```
kubectl completion SHELL [options]
```

```
=====

kubectl alpha --help
```

These commands correspond to alpha features that are not enabled in Kubernetes clusters by default.

Available Commands:

```
auth          Inspect authorization
```

```
=====
```



```
kubectl api-resources --help
Print the supported API resources on the server.
```

Examples:

```
# Print the supported API resources
kubectl api-resources

# Print the supported API resources with more information
kubectl api-resources -o wide

# Print the supported API resources sorted by a column
kubectl api-resources --sort-by=name

# Print the supported namespaced resources
kubectl api-resources --namespaced=true

# Print the supported non-namespaced resources
kubectl api-resources --namespaced=false

# Print the supported API resources with a specific APIGroup
kubectl api-resources --api-group=rbac.authorization.k8s.io
```

Options:

```
--api-group='':
    Limit to resources in the specified API group.

--cached=false:
    Use the cached list of resources if available.

--categories=[ ]:
    Limit to resources that belong the the specified categories.

--namespaced=true:
    If false, non-namespaced resources will be returned, otherwise
    returning namespaced resources by default.

--no-headers=false:
    When using the default or custom-column output format, don't print
    headers (default print headers).

-o, --output='':
    Output format. One of: (wide, name).

--sort-by='':
    If non-empty, sort list of resources using specified field. The field
    can be either 'name' or 'kind'.

--verbs=[ ]:
    Limit to resources that support the specified verbs.
```

Usage:

```
kubectl api-resources [flags] [options]
```

=====

```
kubectl api-versions --help
Print the supported API versions on the server, in the form of "group/version".
```

Examples:

```
# Print the supported API versions
kubectl api-versions
```

Usage:

```
kubectl api-versions [options]
```

```
=====
kubecttl config --help
Modify kubeconfig files using subcommands like "kubecttl config set
current-context my-context"
```

The loading order follows these rules:

1. If the --kubeconfig flag is set, then only that file is loaded. The flag may only be set once and no merging takes place.
2. If \$KUBECONFIG environment variable is set, then it is used as a list of paths (normal path delimiting rules for your system). These paths are merged. When a value is modified, it is modified in the file that defines the stanza. When a value is created, it is created in the first file that exists. If no files in the chain exist, then it creates the last file in the list.
3. Otherwise, \${HOME}/.kube/config is used and no merging takes place.

Available Commands:

current-context	Display the current-context
delete-cluster	Delete the specified cluster from the kubeconfig
delete-context	Delete the specified context from the kubeconfig
delete-user	Delete the specified user from the kubeconfig
get-clusters	Display clusters defined in the kubeconfig
get-contexts	Describe one or many contexts
get-users	Display users defined in the kubeconfig
rename-context	Rename a context from the kubeconfig file
set	Set an individual value in a kubeconfig file
set-cluster	Set a cluster entry in kubeconfig
set-context	Set a context entry in kubeconfig
set-credentials	Set a user entry in kubeconfig
unset	Unset an individual value in a kubeconfig file
use-context	Set the current-context in a kubeconfig file
view	Display merged kubeconfig settings or a specified kubeconfig file

Usage:

```
kubecttl config SUBCOMMAND [options]
```

```
=====
kubecttl plugin --help
Provides utilities for interacting with plugins.
```

Plugins provide extended functionality that is not part of the major command-line distribution. Please refer to the documentation and examples for more information about how write your own plugins.

The easiest way to discover and install plugins is via the kubernetes sub-project krew. To install krew, visit
<https://krew.sigs.k8s.io/docs/user-guide/setup/install/> krew.sigs.k8s.io
<https://krew.sigs.k8s.io/docs/user-guide/setup/install/>

Available Commands:

list	List all visible plugin executables on a user's PATH
------	--

Usage:

```
kubecttl plugin [flags] [options]
```

```
=====
kubecttl version --help
Print the client and server version information for the current context.
```

Examples:

```
# Print the client and server versions for the current context
kubectl version
```

Options:

```
--client=false:
    If true, shows client version only (no server required).

-o, --output='':
    One of 'yaml' or 'json'.
```

Usage:

```
kubectl version [flags] [options]
```

=====

Password:

```
[root@freeipa tm]# cat <<EOF > /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://packages.cloud.google.com/yum/repos/kubernetes-el7-x86_64
enabled=1
gpgcheck=1
repo_gpgcheck=1
gpgkey=https://packages.cloud.google.com/yum/doc/yum-key.gpg
https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
EOF
[root@freeipa tm]# dnf install -y kubelet kubeadm kubectl
[sudo] password for tm:
Extra Packages for Enterprise Linux 9 - x86_64    24 kB/s | 29 kB    00:01
Extra Packages for Enterprise Linux 9 - x86_64  408 kB/s | 15 MB    00:36
Kubernetes                                     177 B/s | 454 B    00:02
Kubernetes                                     11 kB/s | 2.6 kB    00:00
Importing GPG key 0x13EDEF05:
  Userid      : "Rapture Automatic Signing Key (cloud-rapture-signing-key-2022-03-07-08_01_01.pub)"
  Fingerprint: A362 B822 F6DE DC65 2817 EA46 B53D C80D 13ED EF05
  From        : https://packages.cloud.google.com/yum/doc/yum-key.gpg
Importing GPG key 0xDC6315A3:
  Userid      : "Artifact Registry Repository Signer <artifact-registry-repository-signer@google.com>"
  Fingerprint: 35BA A0B3 3E9E B396 F59C A838 C0BA 5CE6 DC63 15A3
  From        : https://packages.cloud.google.com/yum/doc/yum-key.gpg
Kubernetes                                     3.2 kB/s | 975 B    00:00
Importing GPG key 0x3E1BA8D5:
  Userid      : "Google Cloud Packages RPM Signing Key <gc-team@google.com>"
  Fingerprint: 3749 E1BA 95A8 6CE0 5454 6ED2 F09C 394C 3E1B A8D5
  From        : https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg
Kubernetes                                     98 kB/s | 168 kB    00:01
Rocky Linux 9 - BaseOS                         5.7 kB/s | 4.1 kB    00:00
Rocky Linux 9 - BaseOS                         279 kB/s | 1.8 MB    00:06
Rocky Linux 9 - AppStream                       5.8 kB/s | 4.5 kB    00:00
Rocky Linux 9 - AppStream                       275 kB/s | 6.6 MB    00:24
Rocky Linux 9 - Extras                          505 B/s | 2.9 kB    00:05
Rocky Linux 9 - Extras                          12 kB/s | 8.5 kB    00:00
RPM Fusion for EL 9 - Free - Updates            8.7 kB/s | 2.8 kB    00:00
RPM Fusion for EL 9 - Free - Updates           160 kB/s | 245 kB    00:01
RPM Fusion for EL 9 - Nonfree - Updates         34 kB/s | 5.6 kB    00:00
RPM Fusion for EL 9 - Nonfree - Updates        66 kB/s | 63 kB     00:00
```

Dependencies resolved.

Package	Arch	Version	Repository	Size
Installing:				
kubeadm	x86_64	1.26.3-0	kubernetes	10 M
kubect1	x86_64	1.26.3-0	kubernetes	11 M
kubelet	x86_64	1.26.3-0	kubernetes	22 M
Installing dependencies:				
conntrack-tools	x86_64	1.4.5-17.el9_1	appstream	210 k
cri-tools	x86_64	1.26.0-0	kubernetes	8.6 M
kubernetes-cni	x86_64	1.2.0-0	kubernetes	17 M
libnetfilter_cthelper	x86_64	1.0.0-22.el9	appstream	23 k
libnetfilter_cttimeout	x86_64	1.0.0-19.el9	appstream	23 k
libnetfilter_queue	x86_64	1.0.5-1.el9	appstream	28 k
socat	x86_64	1.7.4.1-5.el9	appstream	300 k

Transaction Summary

Install 10 Packages

Total download size: 69 M

Installed size: 296 M

Downloading Packages:

(1/10):	3f5ba2b53701ac9102ea7c7ab2ca6616a8cd596	117 kB/s 8.6 MB	01:15
(2/10):	6dee5ef05a942f42f7e49d2aac919d1177d8006	129 kB/s 10 MB	01:22
(3/10):	8c423fd76dc6d7a5dfc2aa30630cb46b6bd8128	114 kB/s 11 MB	01:35
(4/10):	conntrack-tools-1.4.5-17.el9_1.x86_64.r	130 kB/s 210 kB	00:01
(5/10):	libnetfilter_cttimeout-1.0.0-19.el9.x86	103 kB/s 23 kB	00:00
(6/10):	socat-1.7.4.1-5.el9.x86_64.rpm	178 kB/s 300 kB	00:01
(7/10):	libnetfilter_cthelper-1.0.0-22.el9.x86	77 kB/s 23 kB	00:00
(8/10):	libnetfilter_queue-1.0.5-1.el9.x86_64.r	108 kB/s 28 kB	00:00
(9/10):	0f2a2afd740d476ad77c508847bad1f559afc24	214 kB/s 17 MB	01:20
(10/10):	4bd2c321343cbe55e4e50e6095273a2bca4ecd	233 kB/s 22 MB	01:35

Total	411 kB/s 69 MB	02:51
Kubernetes	8.9 kB/s 2.6 kB	00:00

Importing GPG key 0x13EDEF05:

 Userid : "Rapture Automatic Signing Key (cloud-rapture-signing-key-2022-03-07-08_01_01.pub)"

 Fingerprint: A362 B822 F6DE DC65 2817 EA46 B53D C80D 13ED EF05

 From : https://packages.cloud.google.com/yum/doc/yum-key.gpg

Key imported successfully

Importing GPG key 0xDC6315A3:

 Userid : "Artifact Registry Repository Signer <artifact-registry-repository-signer@google.com>"

 Fingerprint: 35BA A0B3 3E9E B396 F59C A838 C0BA 5CE6 DC63 15A3

 From : https://packages.cloud.google.com/yum/doc/yum-key.gpg

Key imported successfully

Kubernetes	3.9 kB/s 975 B	00:00
------------	------------------	-------

Importing GPG key 0x3E1BA8D5:

 Userid : "Google Cloud Packages RPM Signing Key <gc-team@google.com>"

 Fingerprint: 3749 E1BA 95A8 6CE0 5454 6ED2 F09C 394C 3E1B A8D5

 From : https://packages.cloud.google.com/yum/doc/rpm-package-key.gpg

Key imported successfully

Running transaction check

Transaction check succeeded.

Running transaction test

Transaction test succeeded.

Running transaction

Preparing	:	1/1
Installing	: libnetfilter_queue-1.0.5-1.el9.x86_64	1/10
Installing	: libnetfilter_cthelper-1.0.0-22.el9.x86_64	2/10
Installing	: socat-1.7.4.1-5.el9.x86_64	3/10
Installing	: libnetfilter_cttimeout-1.0.0-19.el9.x86_64	4/10
Installing	: conntrack-tools-1.4.5-17.el9_1.x86_64	5/10

Running scriptlet:	conntrack-tools-1.4.5-17.el9_1.x86_64	5/10
Installing	: kubernetes-cni-1.2.0-0.x86_64	6/10
Installing	: kubelet-1.26.3-0.x86_64	7/10
Installing	: kubect1-1.26.3-0.x86_64	8/10
Installing	: cri-tools-1.26.0-0.x86_64	9/10
Installing	: kubeadm-1.26.3-0.x86_64	10/10
Running scriptlet:	kubeadm-1.26.3-0.x86_64	10/10
Verifying	: cri-tools-1.26.0-0.x86_64	1/10
Verifying	: kubeadm-1.26.3-0.x86_64	2/10
Verifying	: kubect1-1.26.3-0.x86_64	3/10
Verifying	: kubelet-1.26.3-0.x86_64	4/10
Verifying	: kubernetes-cni-1.2.0-0.x86_64	5/10
Verifying	: conntrack-tools-1.4.5-17.el9_1.x86_64	6/10
Verifying	: libnetfilter_cttimeout-1.0.0-19.el9.x86_64	7/10
Verifying	: socat-1.7.4.1-5.el9.x86_64	8/10
Verifying	: libnetfilter_cthelper-1.0.0-22.el9.x86_64	9/10
Verifying	: libnetfilter_queue-1.0.5-1.el9.x86_64	10/10

Installed:

```

conntrack-tools-1.4.5-17.el9_1.x86_64
cri-tools-1.26.0-0.x86_64
kubeadm-1.26.3-0.x86_64
kubect1-1.26.3-0.x86_64
kubelet-1.26.3-0.x86_64
kubernetes-cni-1.2.0-0.x86_64
libnetfilter_cthelper-1.0.0-22.el9.x86_64
libnetfilter_cttimeout-1.0.0-19.el9.x86_64
libnetfilter_queue-1.0.5-1.el9.x86_64
socat-1.7.4.1-5.el9.x86_64

```